

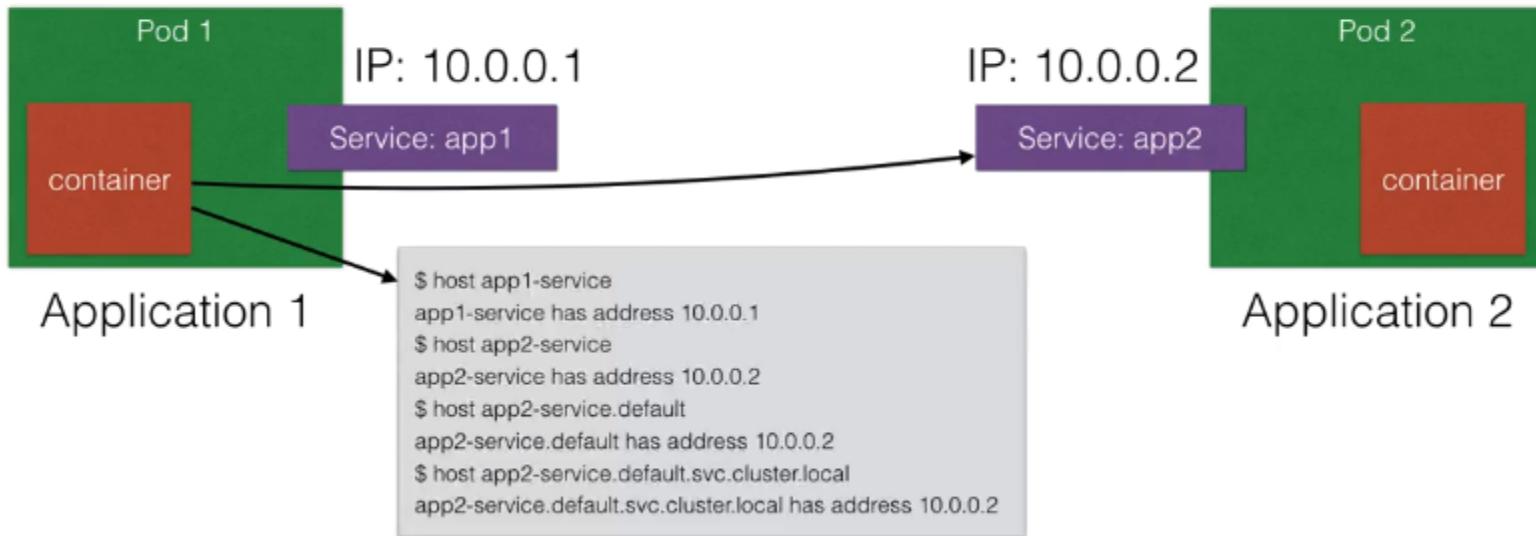
# Service Discovery

# DNS

- As of Kubernetes 1.10, DNS is a **built-in** service launched automatically using addon manger
- The DNS Service can be used within pods to find other service running on the same cluster
- Multiple containers **within 1 pod** don't need this service, as they can contact each other **directly**
- A container in the same pod can **connect / communicate** to other containers directly using **localhost:port**

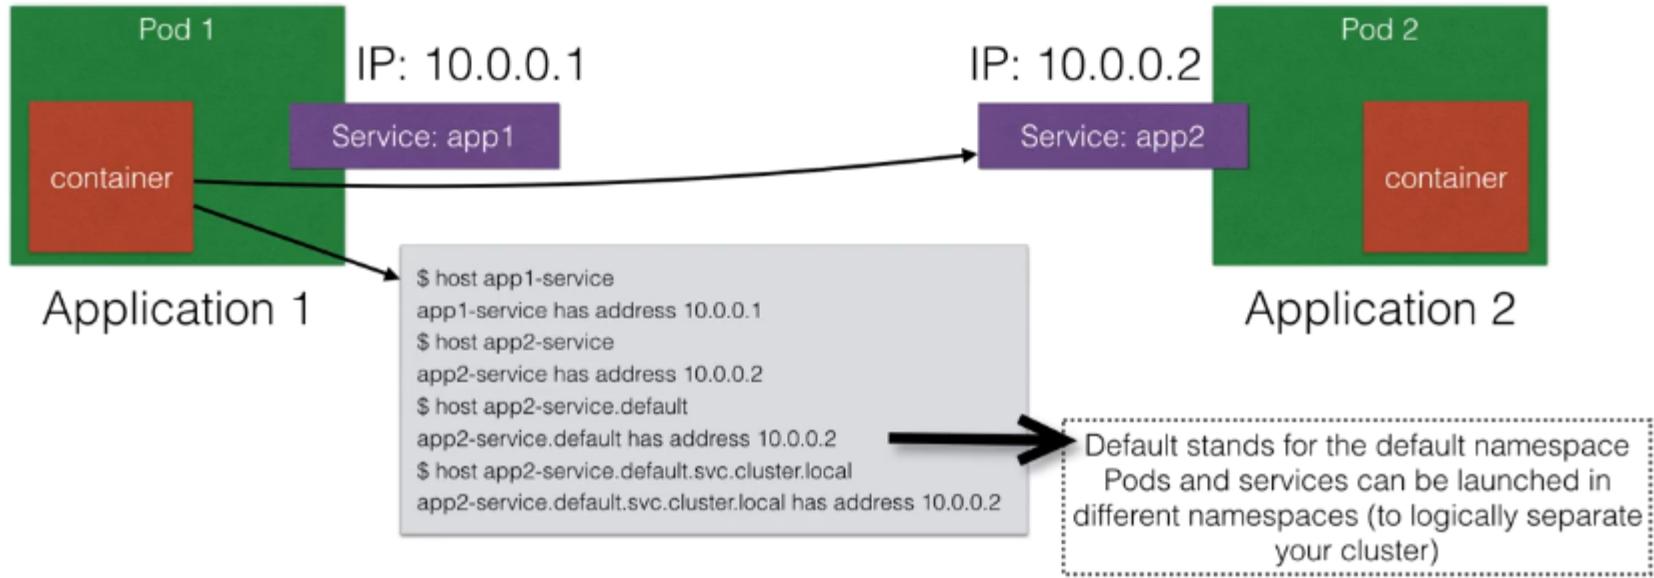
# DNS

- An example of how app 1 could reach app 2 using DNS:

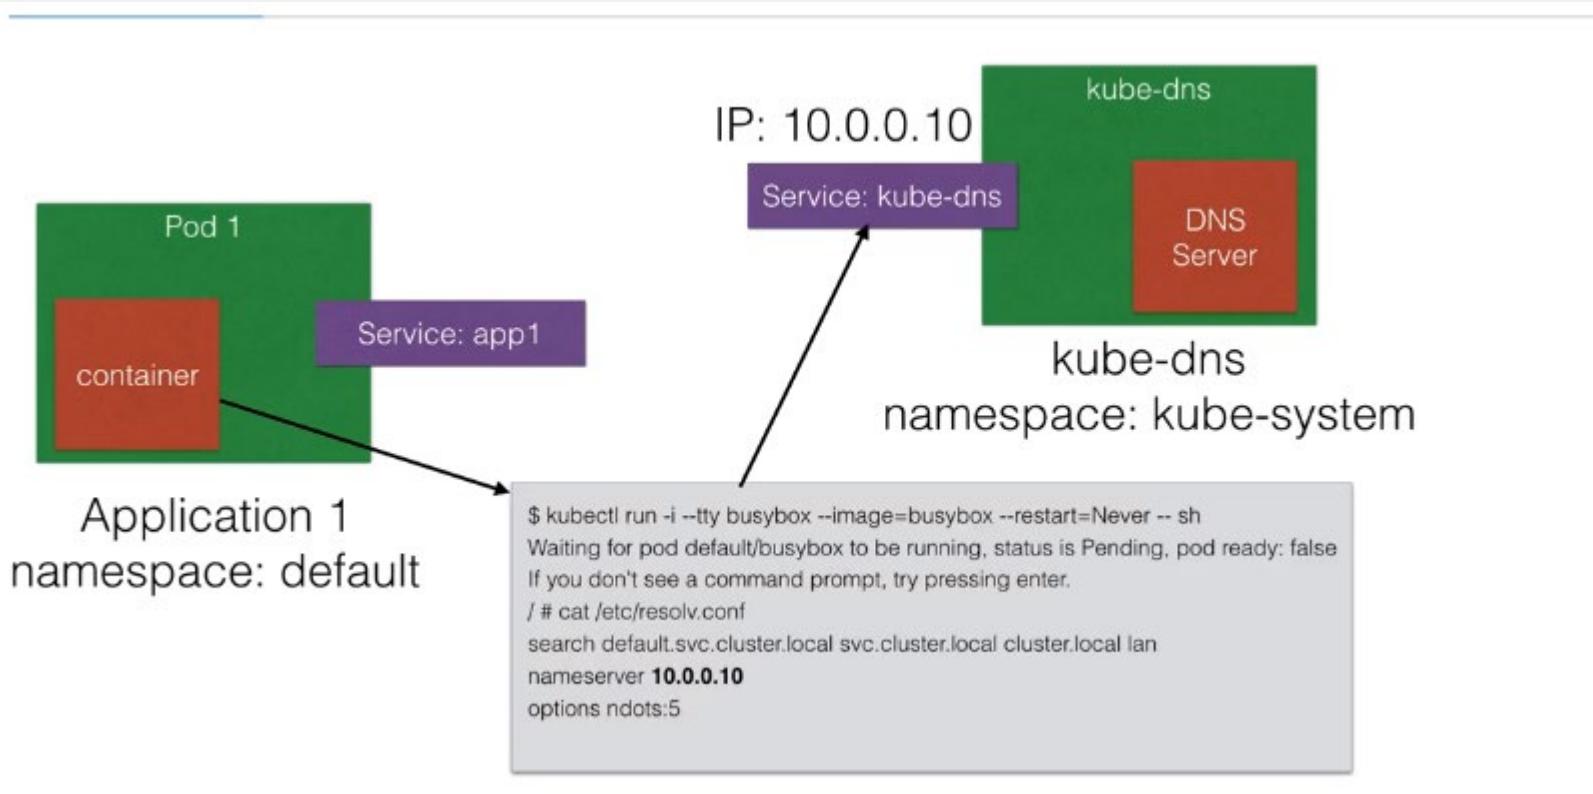


# DNS

- An example of how app 1 could reach app 2 using DNS:



# DNS – How does it work ?



# Demo Placeholder

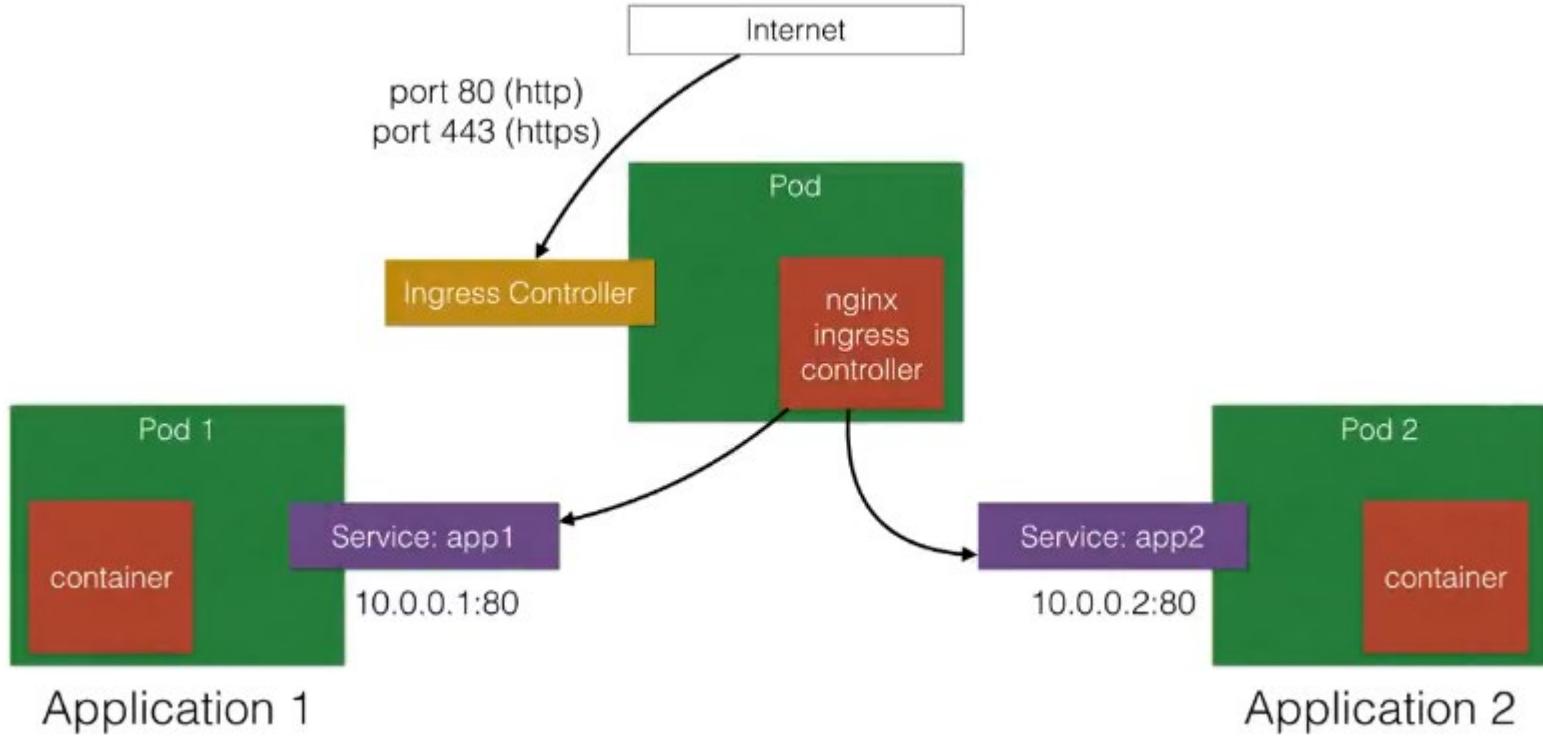
Service Discovery

# Ingress

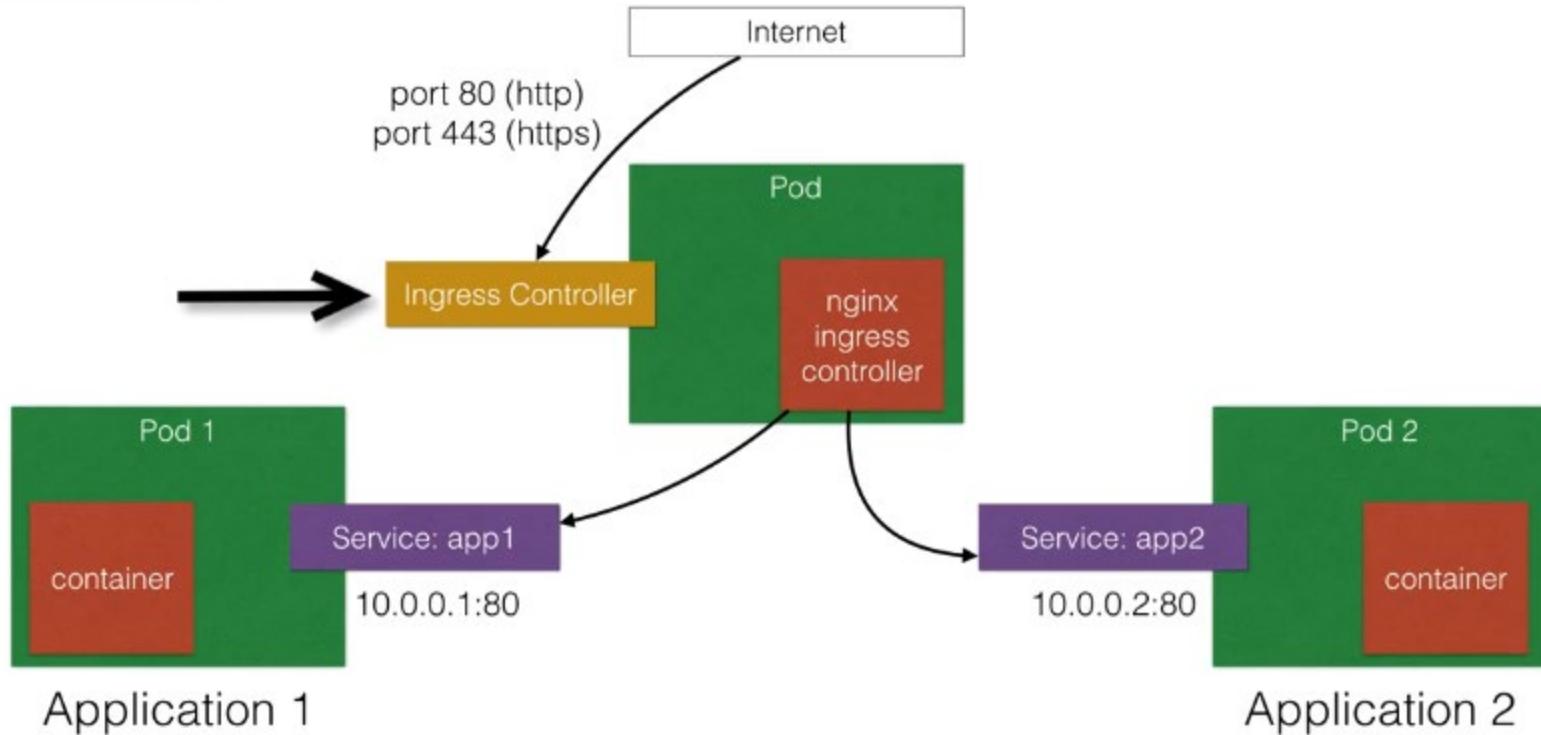
# Ingress

- Ingress is a solution available since kubernetes 1.1 that allows **inbound connection** to the cluster
- It's an alternative to the external LoadBalancer & NodePorts
  - Ingress allows you to **easily expose services** that need to be accessible from **outside** to the **cluster**
- With ingress you can run your own **ingress controller** ( basically a LoadBalancer ) within the kubernetes cluster
- There are a default ingress controllers available, or you can **write your own** ingress controller

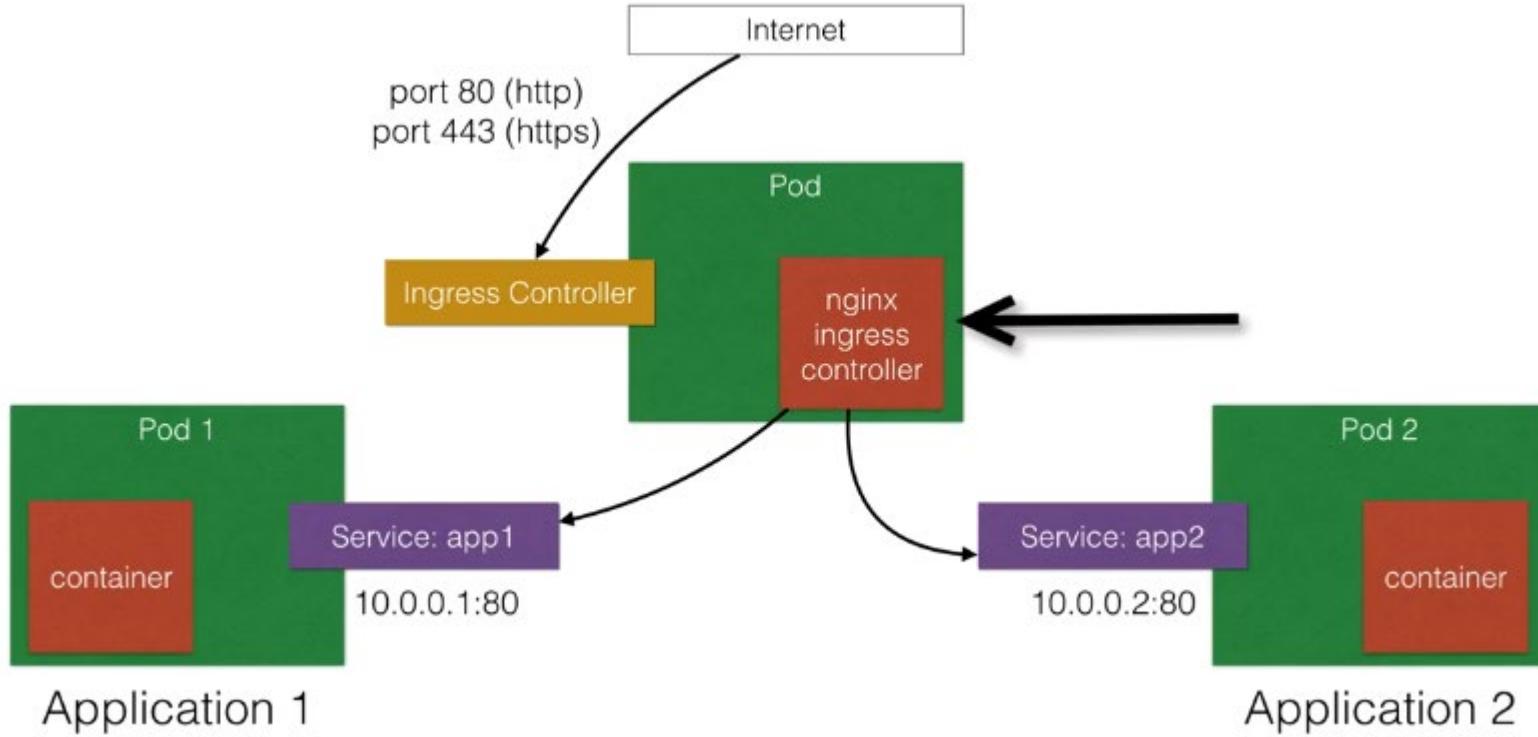
# Ingress



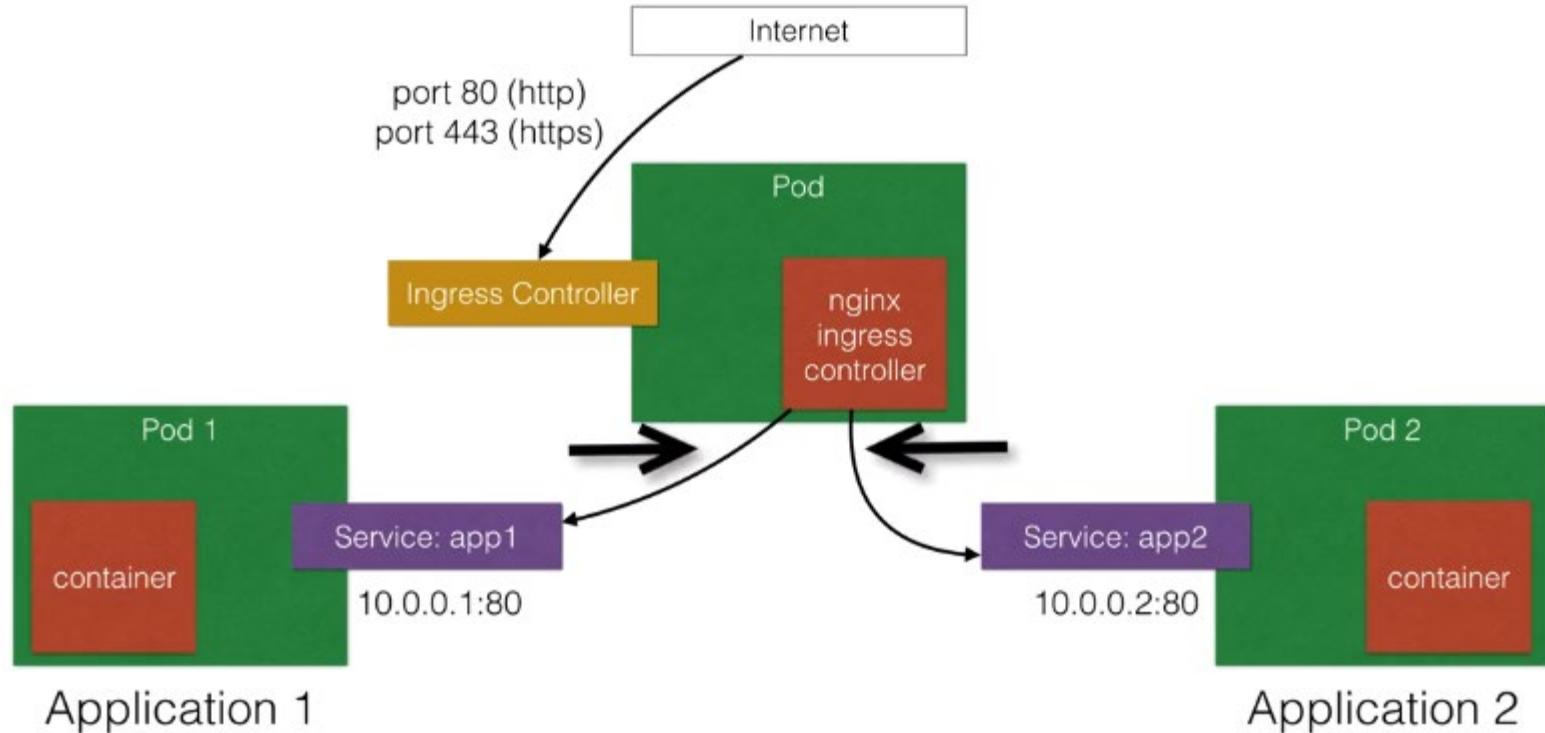
# Ingress



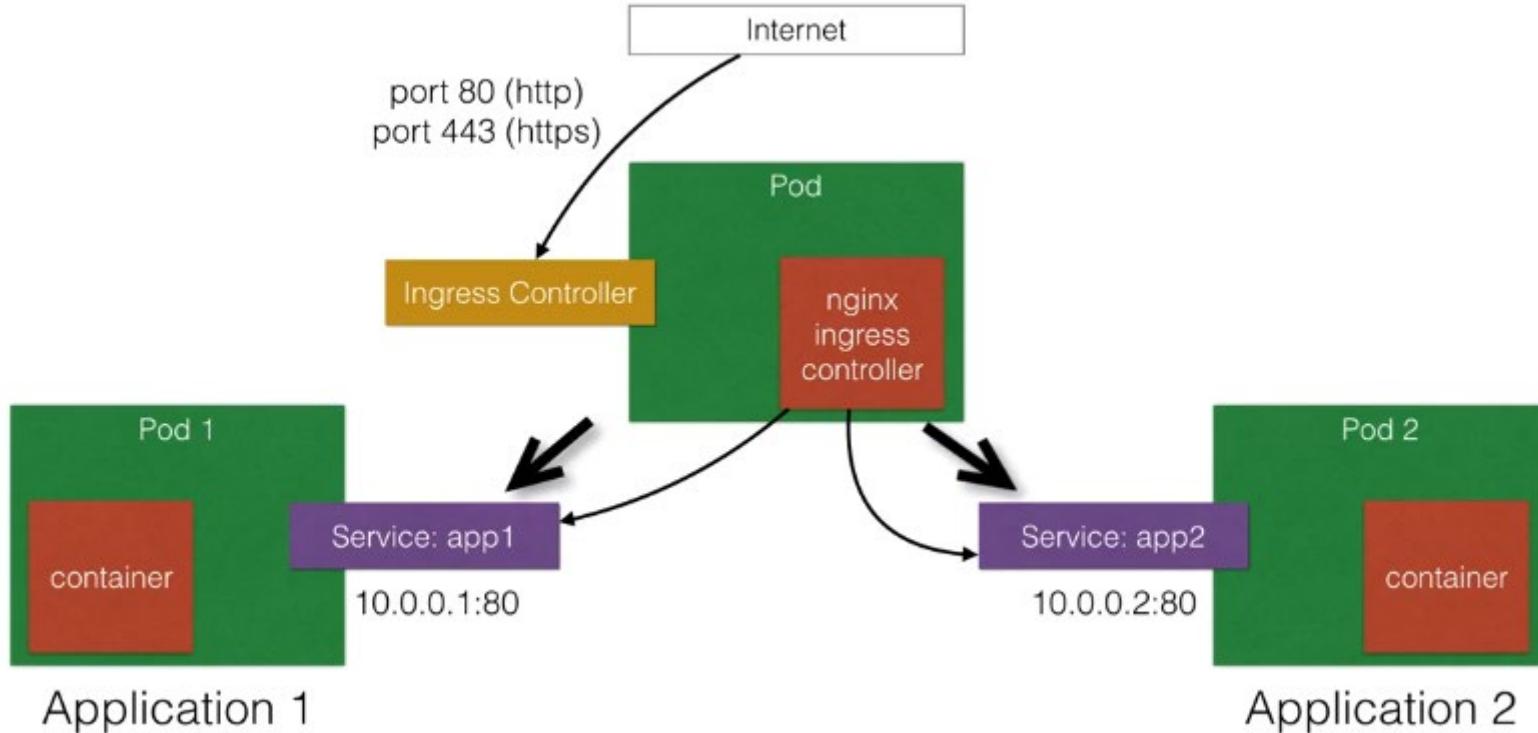
# Ingress



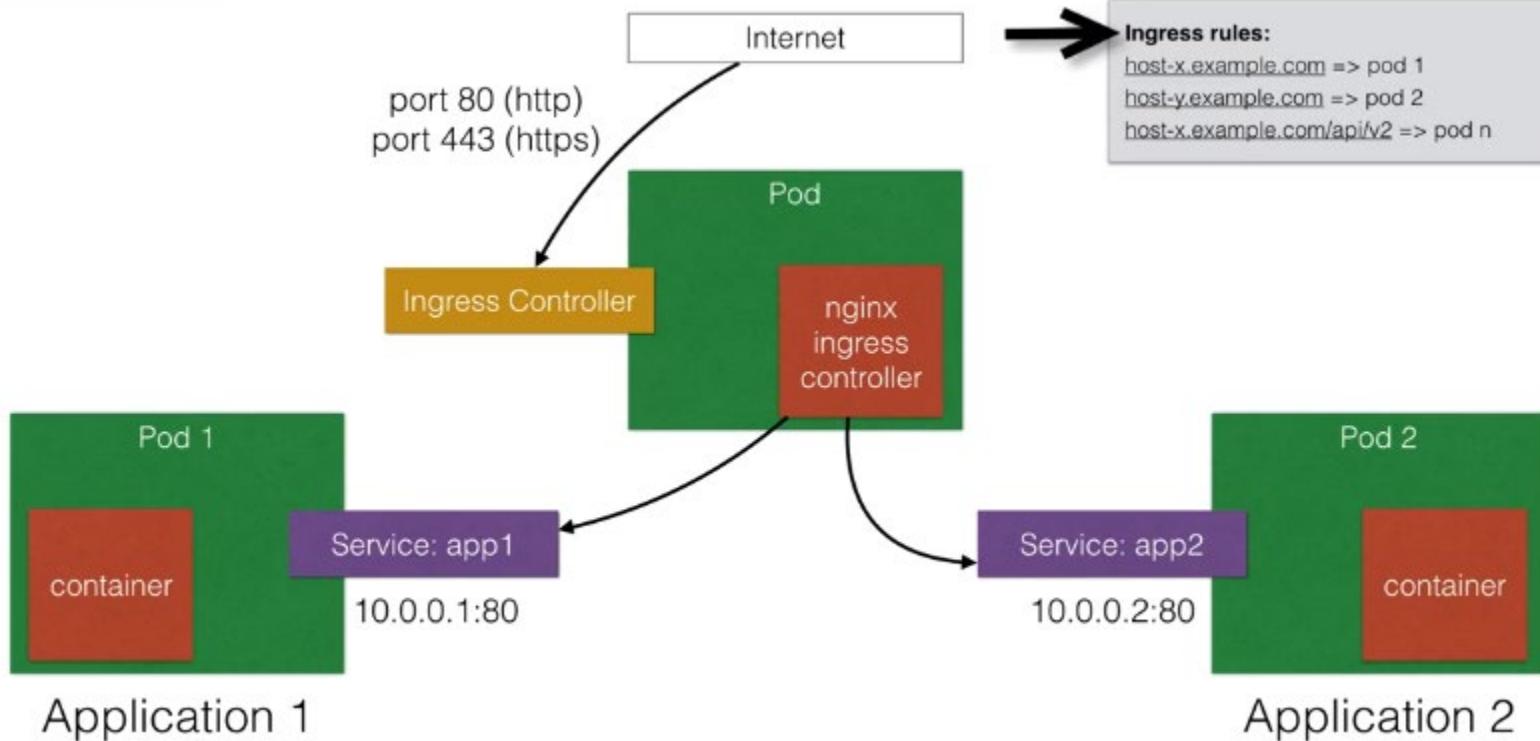
# Ingress



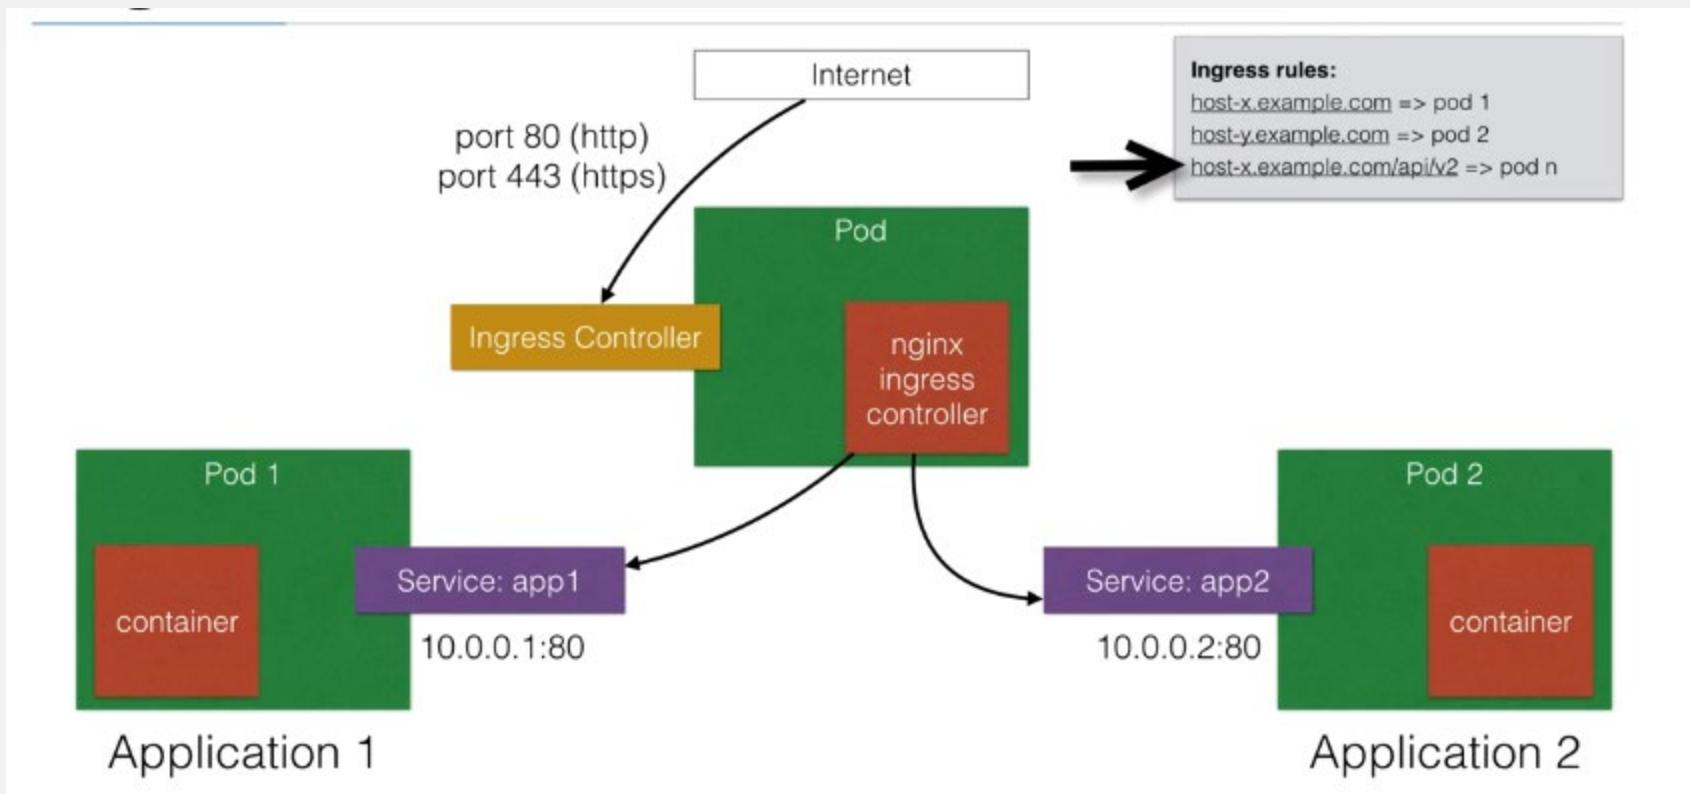
# Ingress



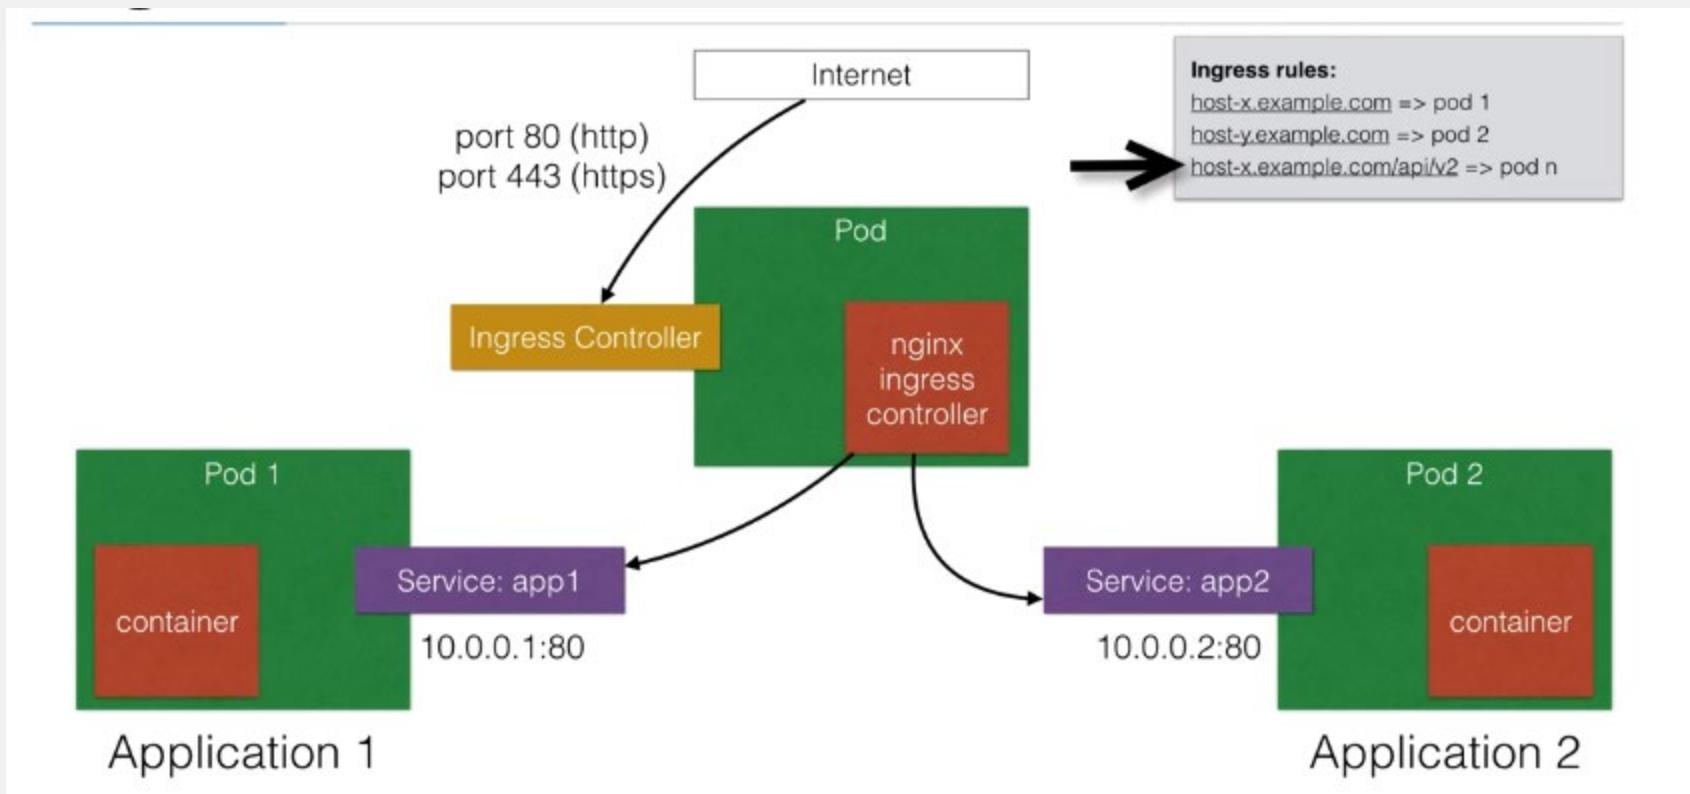
# Ingress



# Ingress



# Ingress



# Ingress

- You can create ingress rules using the ingress objects

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: helloworld-rules
spec:
  rules:
    - host: helloworld-v1.example.com
      http:
        paths:
          - path: /
            backend:
              serviceName: helloworld-v1
              servicePort: 80
    - host: helloworld-v2.example.com
      http:
        paths:
          - path: /
            backend:
              serviceName: helloworld-v2
              servicePort: 80
```



# Demo Placeholder

Ingress Controller

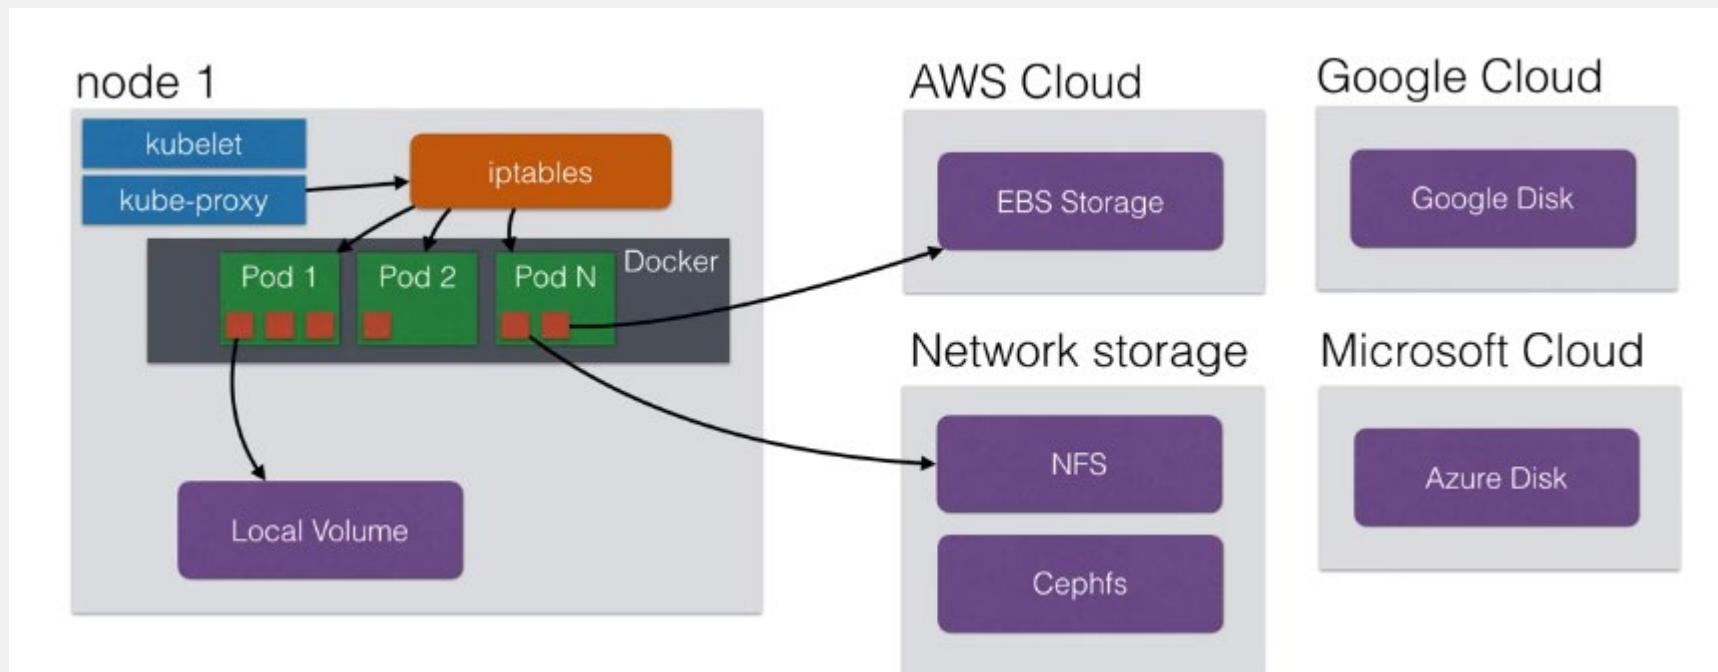
# Volumes

# Volumes

- Volumes in kubernetes allow you to **store data outside the container**
- When a containers **stops**, all data on the container itself is **lost**
  - That's why up until now I've been using **stateless** apps: apps that don't keep a **local** state, but store theirs state in an **external service**
    - External Service like a database, cashing server ( e.g. MySQL, AWS S3 )
- Persistent Volumes in Kubernetes allow you **attach a volume** to a container that will **exists** even when the **container** stops

# Kubernetes Volumes

- Volumes can be attached using different volumes plugins:

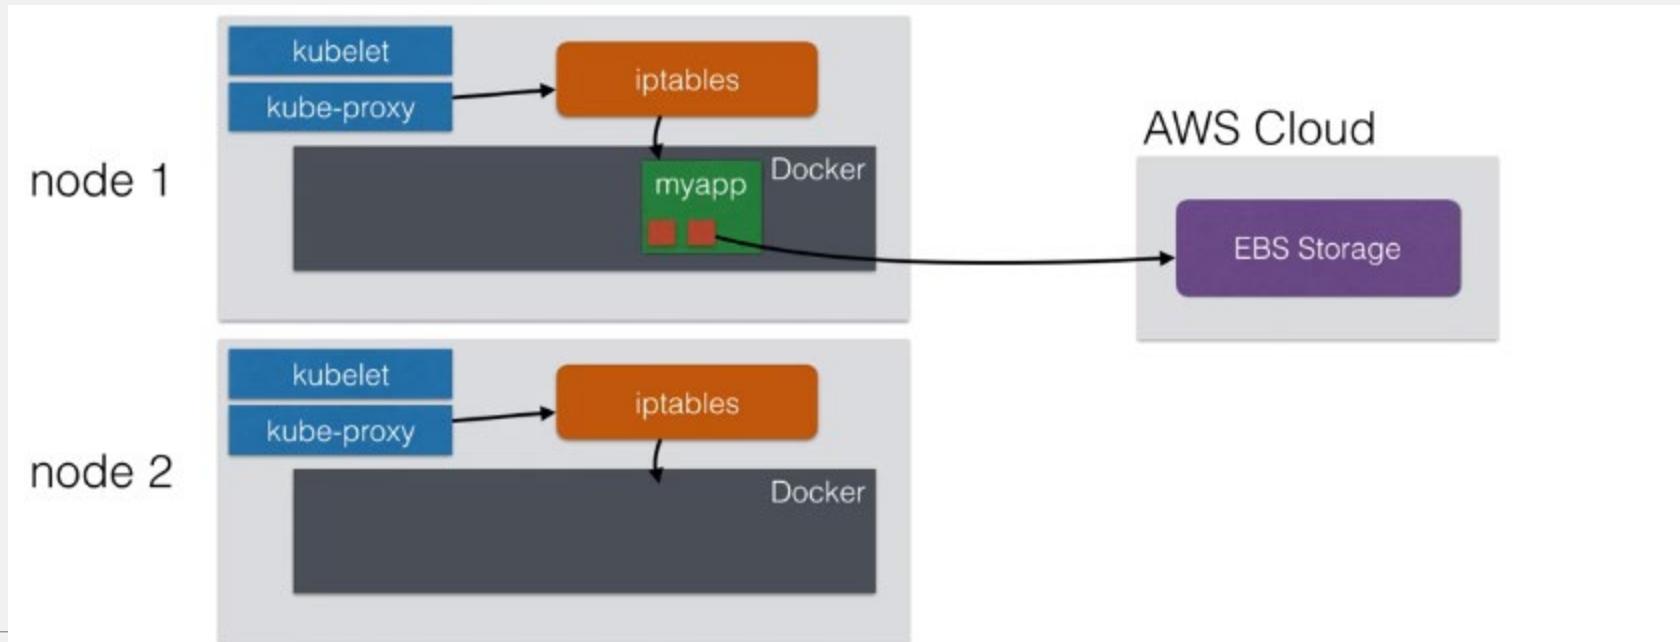


# Volumes

- Using volumes, you could deploy **applications with state** on your cluster
  - Those applications need to read/write to files on the **local filesystem** that need to be persistent in time
- You could run a MySQL database using persistent volumes
  - Although this might not be ready for production ( yet )
  - Volumes are new since the June 2016 release in Kubernetes

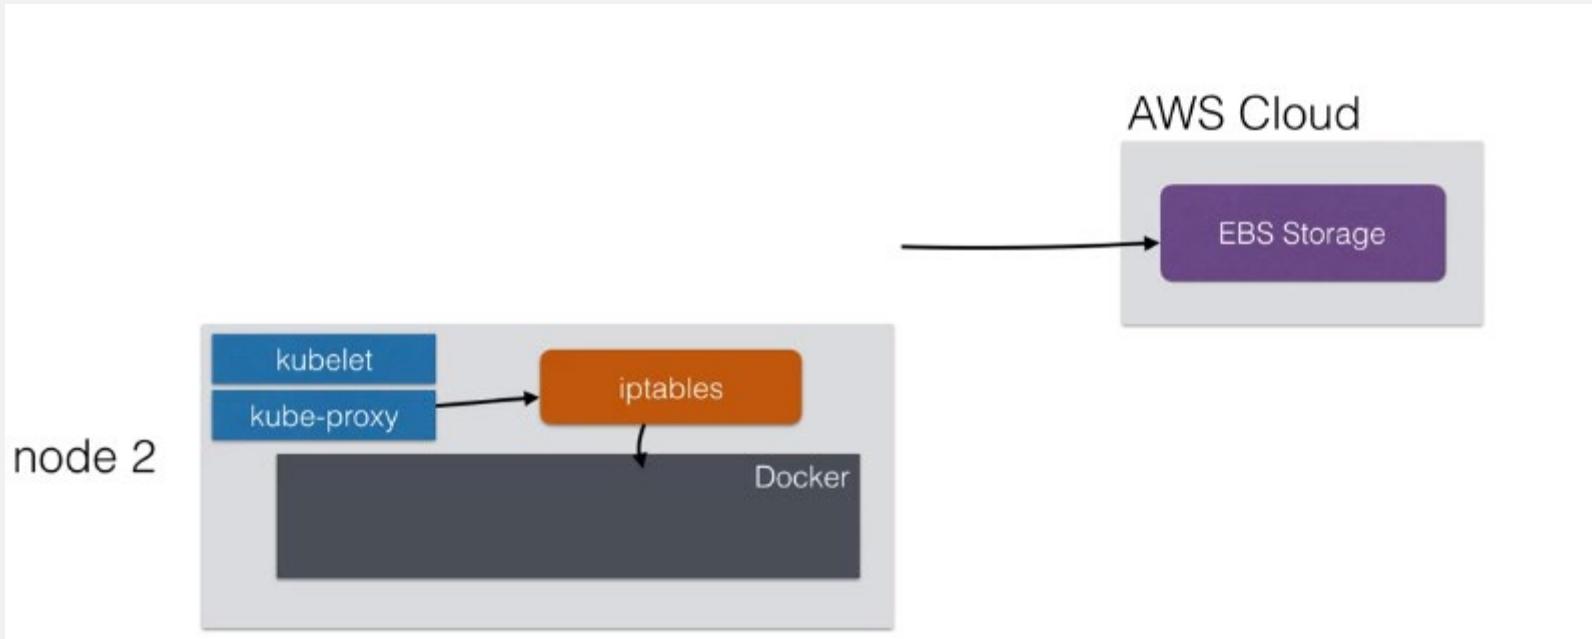
# Volumes

- If your **node** stops working, the pod can be rescheduled on another node, and the volumes can be attached to new node



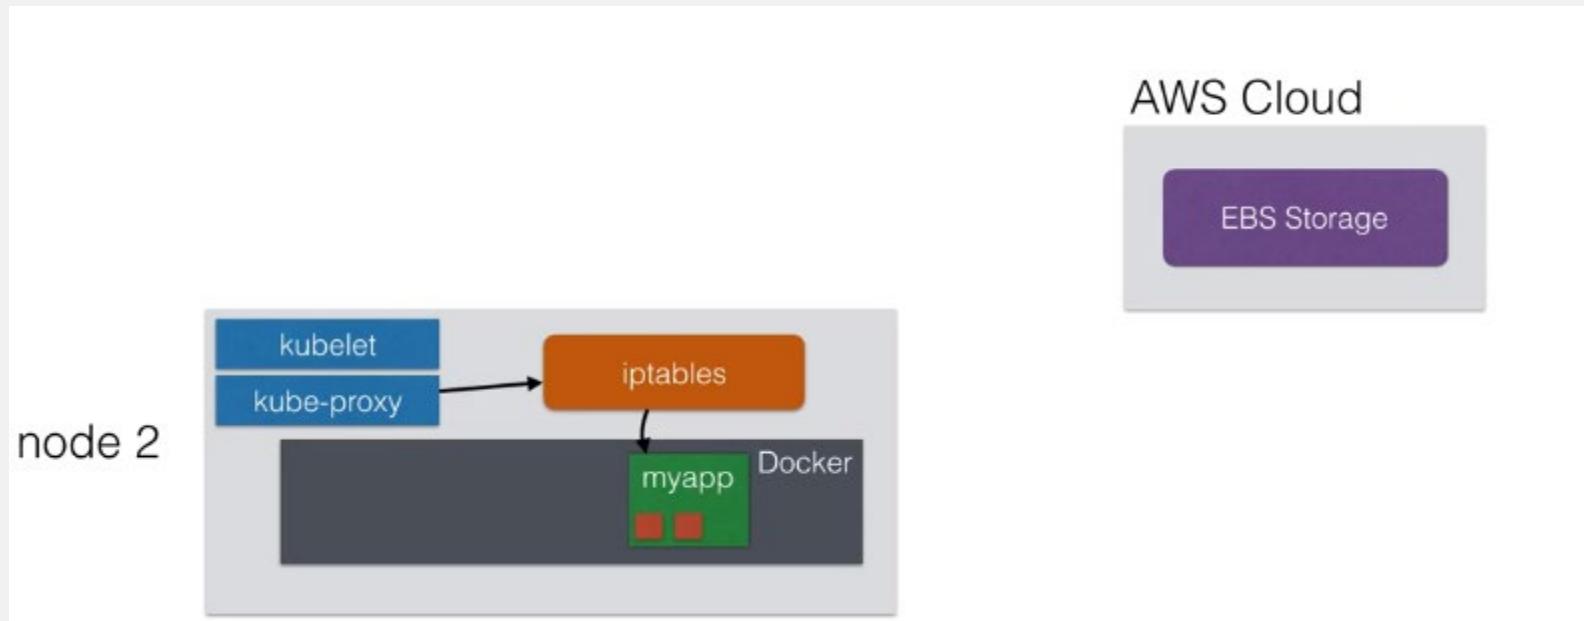
# Volumes

- If your **node** stops working, the pod can be rescheduled on another node, and the volumes can be attached to new node



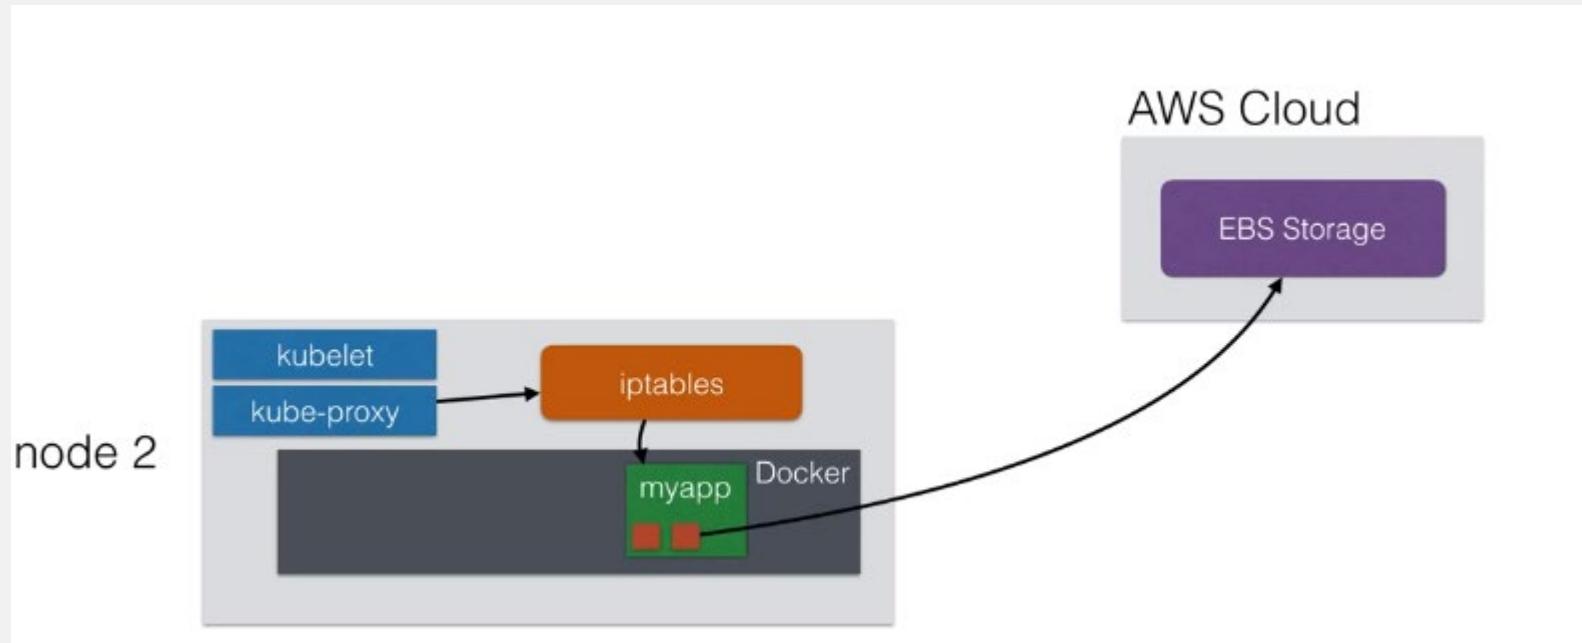
# Volumes

- If your **node** stops working, the pod can be rescheduled on another node, and the volumes can be attached to new node



# Volumes

- If your **node** stops working, the pod can be rescheduled on another node, and the volumes can be attached to new node



# Pod State

# Pod State

- Pods have a status field, which you see when you do *kubectl get pods*:

```
$ kubectl get pods -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE
dns-controller-7cc97fb976-4b9nt   1/1     Running   0          4h
etcd-server-events-ip-172-20-38-169.eu-west-1.compute.internal   1/1     Running   0          4h
etcd-server-ip-172-20-38-169.eu-west-1.compute.internal        1/1     Running   0          4h
kube-apiserver-ip-172-20-38-169.eu-west-1.compute.internal    1/1     Running   0          4h
```

- In this scenario all pods are in the **running status**
  - This means that the **pod has been bound to a node**
  - All **containers have been created**
  - **At least one container** is still **running**, or is starting/restarting

# Pod State

- Other valid statuses are:
  - **Pending**: Pod has been **accepted** but is **not running**
    - Happens when the container image is still **downloading**
    - If the pod cannot be scheduled because of **resource constraints**, it'll also be in this status
  - **Succeeded**: All containers within this pod have been **terminated successfully** and will not be restarted

# Pod State

- Other valid statuses are:
  - **Failed:** All containers within this pod have been **Terminated**, and at least one container returned a failure code
    - The failure code is the **exit code** of the process when a container terminates
  - **Unknown:** The **state of the pod couldn't be determined**
    - A **network error** might have occurred (for example the node where the pod is running on is down)

# Pod State

- You can get the pod conditions using kubectl describe pod PODNAME

```
$ kubectl describe pod kube-apiserver-ip-172-20-38-169.eu-west-1.compute.internal -n kube-system
[...]
Conditions:
  Type      Status
  Initialized  True
  Ready      True
  PodScheduled  True
```

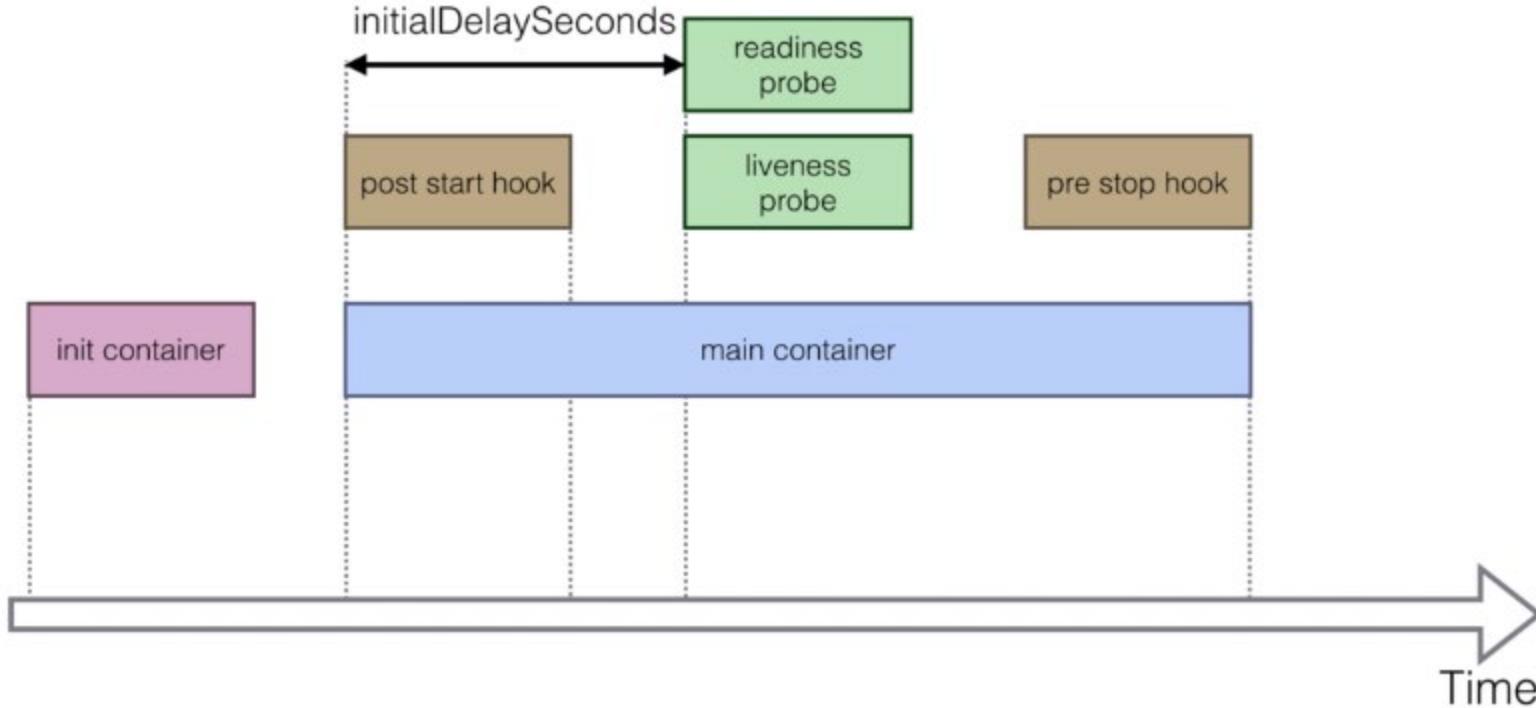
- These are conditions which the pod has passed
  - In this example, Initialized, Ready, and PodScheduled

# Pod State

- There are 5 different types of PodConditions:
  - **PodScheduled**: the pod has been scheduled to a node
  - **Ready**: Pod can serve requests and is going to be added to matching Services
  - **Initialized**: the initialization containers have been started successfully
  - **Unschedulable**: the Pod can't be scheduled (for example due to resource constraints)
  - **ContainersReady**: all containers in the pod are ready

# Pod LifeCycle

# Pod Lifecycle



# Demo

# StatefulSets

# Stateful Sets

- Pet Sets was a **new feature** starting from Kubernetes 1.3, and got renamed to StatefulSets which is stable since Kubernetes 1.9
- It is introduced to be able to run **stateful applications**:
  - That need a **stable pod hostname** (instead of podname-randomstring)
  - Your podname will have a sticky identity, using an index, e.g. podname-0 podname-1 and podname-2 (and when a pod gets rescheduled, it'll keep that identity)
  - Statefulsets allow **stateful apps stable storage** with volumes based on their ordinal number (podname-**x**)
  - **Deleting** and/or **scaling** a **StatefulSet down** will not delete the volumes associated with the StatefulSet (preserving data)

# Stateful Sets

- A StatefulSet will allow your stateful app to use **DNS** to find other **peers**
  - Cassandra clusters, ElasticSearch clusters, use **DNS** to find other members of the cluster
    - for example: **cassandra-0.cassandra** for all pods to reach the first node in the cassandra cluster
  - Using StatefulSet you can run for instance 3 cassandra nodes on Kubernetes named cassandra-0 until cassandra-2
  - If you wouldn't use StatefulSet, you would get a dynamic hostname, which you wouldn't be able to use in your configuration files, as the name can always change

# Stateful Sets

- A StatefulSet will also allow your stateful app to **order the startup and teardown**:
  - Instead of randomly terminating one pod (one instance of your app), you'll know which one that will go
    - When **scaling up** it goes from 0 to n-1 ( $n = \text{replication factor}$ )
    - When **scaling down** it starts with the highest number ( $n-1$ ) to 0
  - This is useful if you first need to **drain** the data from a node before it can be shut down

# Stateful Sets

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: helloworld-statefull
  labels:
    app: helloworld-statefull
spec:
  serviceName: helloworld-statefull
  replicas: 3
  selector:
    matchLabels:
      app: helloworld-statefull
  template:
    metadata:
      labels:
        app: helloworld-statefull
    spec:
      containers:
        - name: k8s-demo
          image: amitvashist7/k8s-tiny-web
          ports:
            - name: node-port
              containerPort: 80
```

# Deamon Sets

# Daemon Sets

- Daemon Sets ensure that **every single node** in the Kubernetes cluster runs the same pod resource
  - This is useful if you want to **ensure** that a certain pod is running on every single kubernetes node
- When a node is **added** to the cluster, a new pod will be **started** automatically
- Same when a node is **removed**, the pod will not be **rescheduled** on another node



# Packaging and Deploying on Kubernetes

# Helm

# Helm

---

- Helm the best way to find, share and use software built for Kubernetes  
(definition from <https://helm.sh/>)
- Helm is a **package manager** for Kubernetes
- It helps you to manage Kubernetes **applications**
- Helm is maintained by the **CNCF - The Cloud Native Computing Foundation** (together with Kubernetes, fluentd, linkerd, and others)
  - It is now maintained in collaboration with **Microsoft, Google, Bitnami** and the **helm contributor community**

# Helm

---

- To start using helm, you first need to download the **helm client**
- You need to run “helm init” to **initialize helm** on the Kubernetes cluster
  - This will install **Tiller**
  - If you have **RBAC installed** (recent clusters have it enabled now by default), you’ll also need add a **ServiceAccount and RBAC rules**
- After this, helm is ready for use, and you can **start installing charts**

# Helm

---

- Helm uses a packaging format called **charts**
  - A **chart** is a **collection** of **files** that **describe** a set of Kubernetes **resources**
  - A single chart can **deploy an app**, a piece of software, or a database
  - It can have dependencies, e.g. to install wordpress chart, you need a mysql chart
  - You can write **your own chart** to deploy your application on Kubernetes using helm

# Helm

- Charts use **templates** that are typically developed by a package maintainer
- They will generate **yaml** files that Kubernetes understands
- You can think of templates as dynamic yaml files, which can contain logic and variables

# Helm

- This is an example of a **template within a chart**:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: {{ .Release.Name }}-configmap
data:
  myvalue: "Hello World"
  drink: {{ .Values.favoriteDrink }}
```

- The favoriteDrink value can then be overridden by the user when running helm install

# Helm – Common Commands

Command	Description
helm init	Install tiller on the cluster
helm reset	Remove tiller from the cluster
helm install	Install a helm chart
helm search	search for a chart
helm list	list releases (installed charts)
helm upgrade	upgrade a release
helm rollback	rollback a release to the previous version

# Demo Placeholder

Let's set helm.



# Helm

## Create your own charts

# Helm Charts

- You can **create helm charts** to **deploy your own apps**
- It's the **recommended** way to deploy your applications on Kubernetes
  - Packaging the app, allows you **deploy the app in 1 command** (instead of using kubectl create / apply)
  - Helm allows for **upgrades** and **rollbacks**
  - Your helm chart is **version controlled**

# Helm Charts

---

- To create the files necessary for a new chart, you can enter the command:

```
helm create mychart
```

# Helm Charts

---

- To create the files necessary for a new chart, you can enter the command:

```
helm create mychart
```

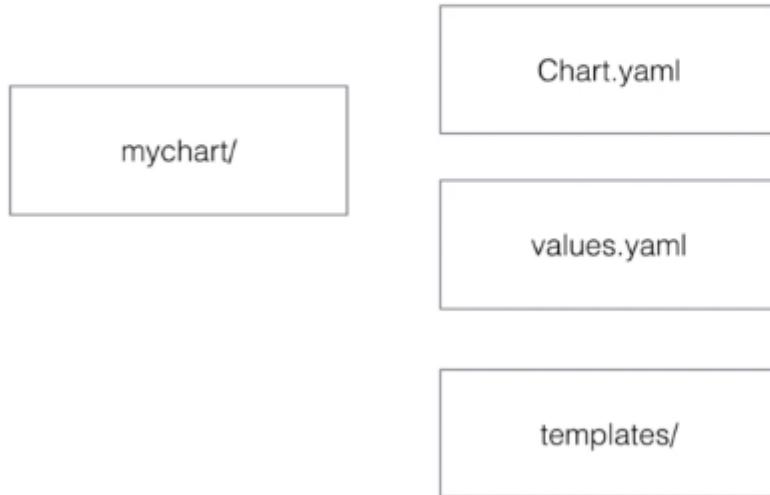
```
mychart/
```

# Helm Charts

---

- To create the files necessary for a new chart, you can enter the command:

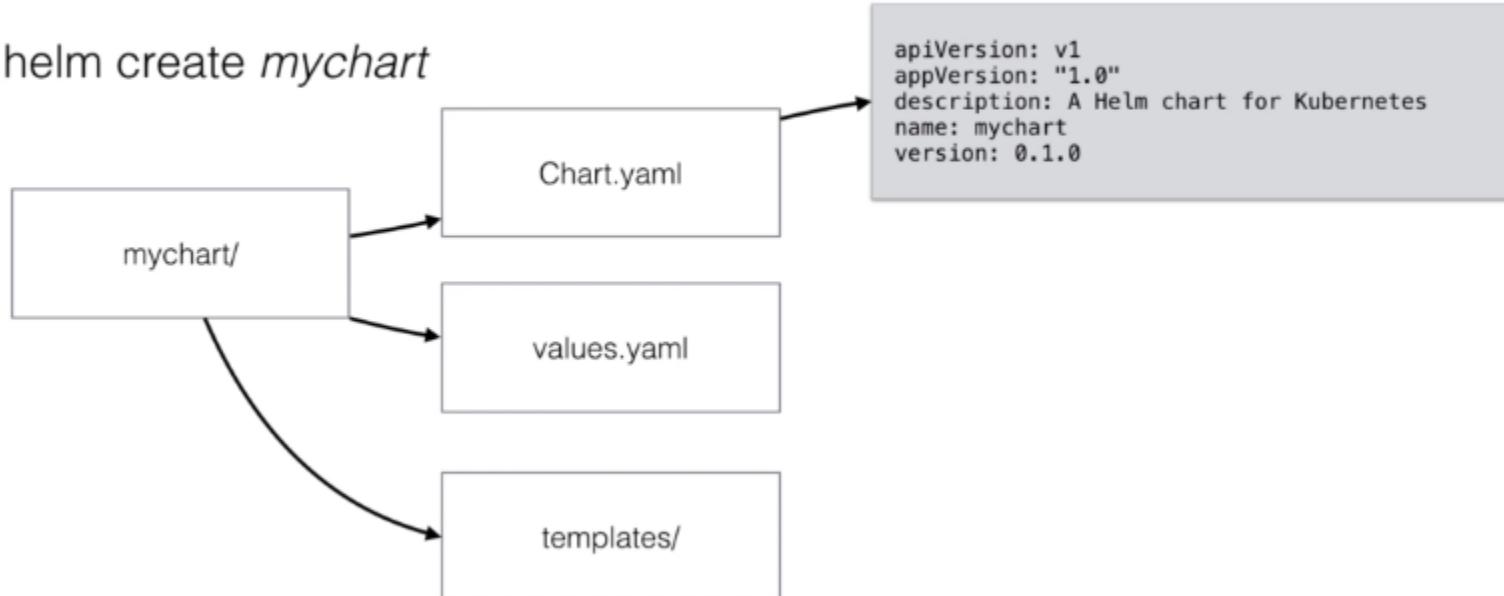
```
helm create mychart
```



# Helm Charts

- To create the files necessary for a new chart, you can enter the command:

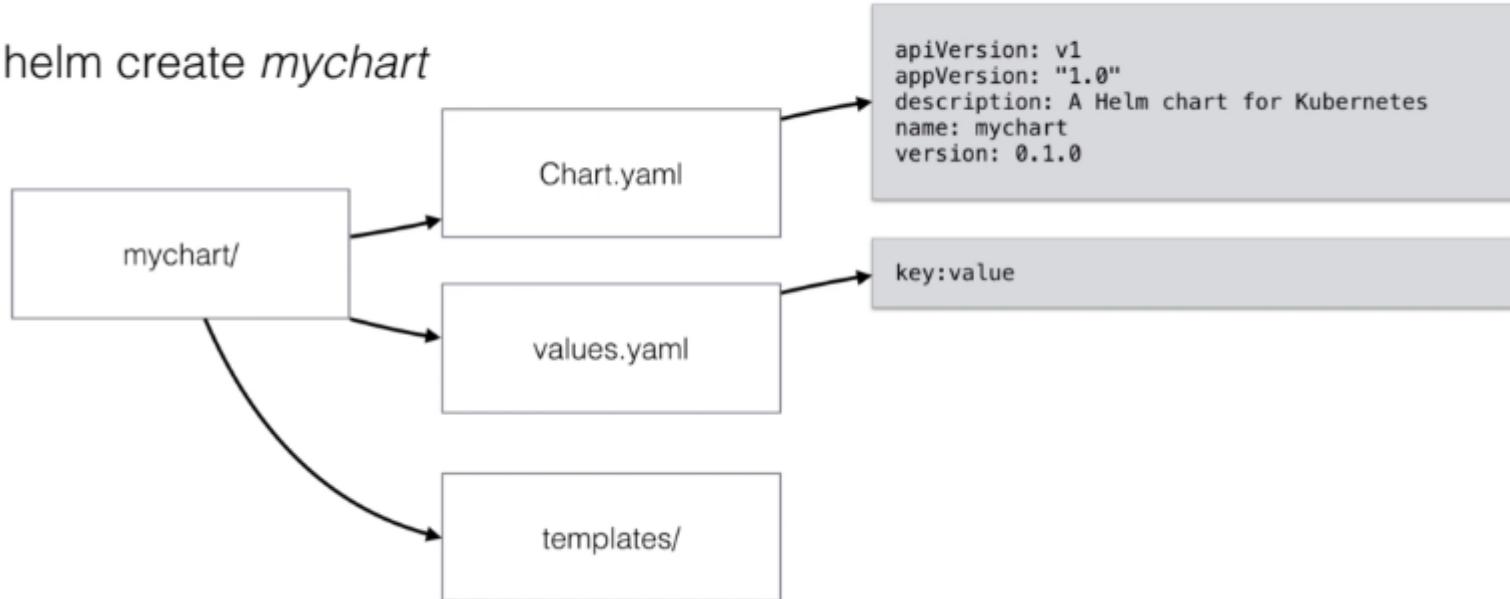
```
helm create mychart
```



# Helm Charts

- To create the files necessary for a new chart, you can enter the command:

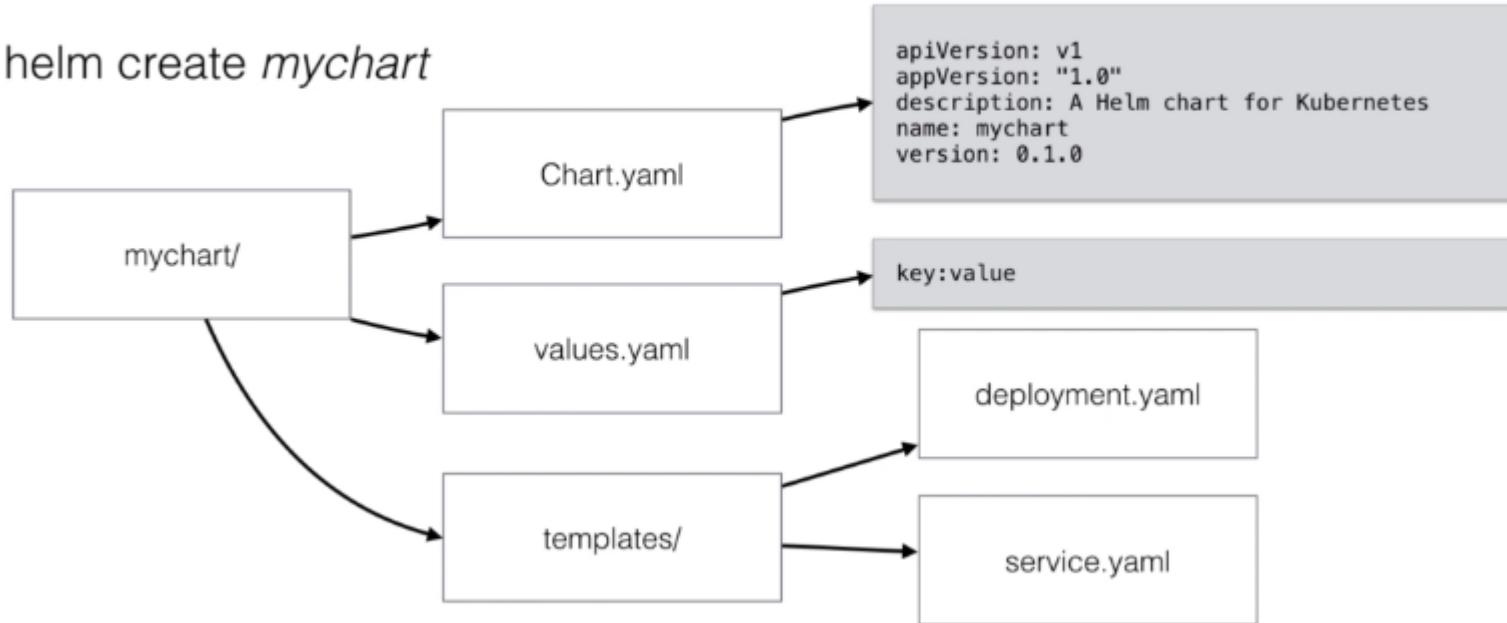
```
helm create mychart
```



# Helm Charts

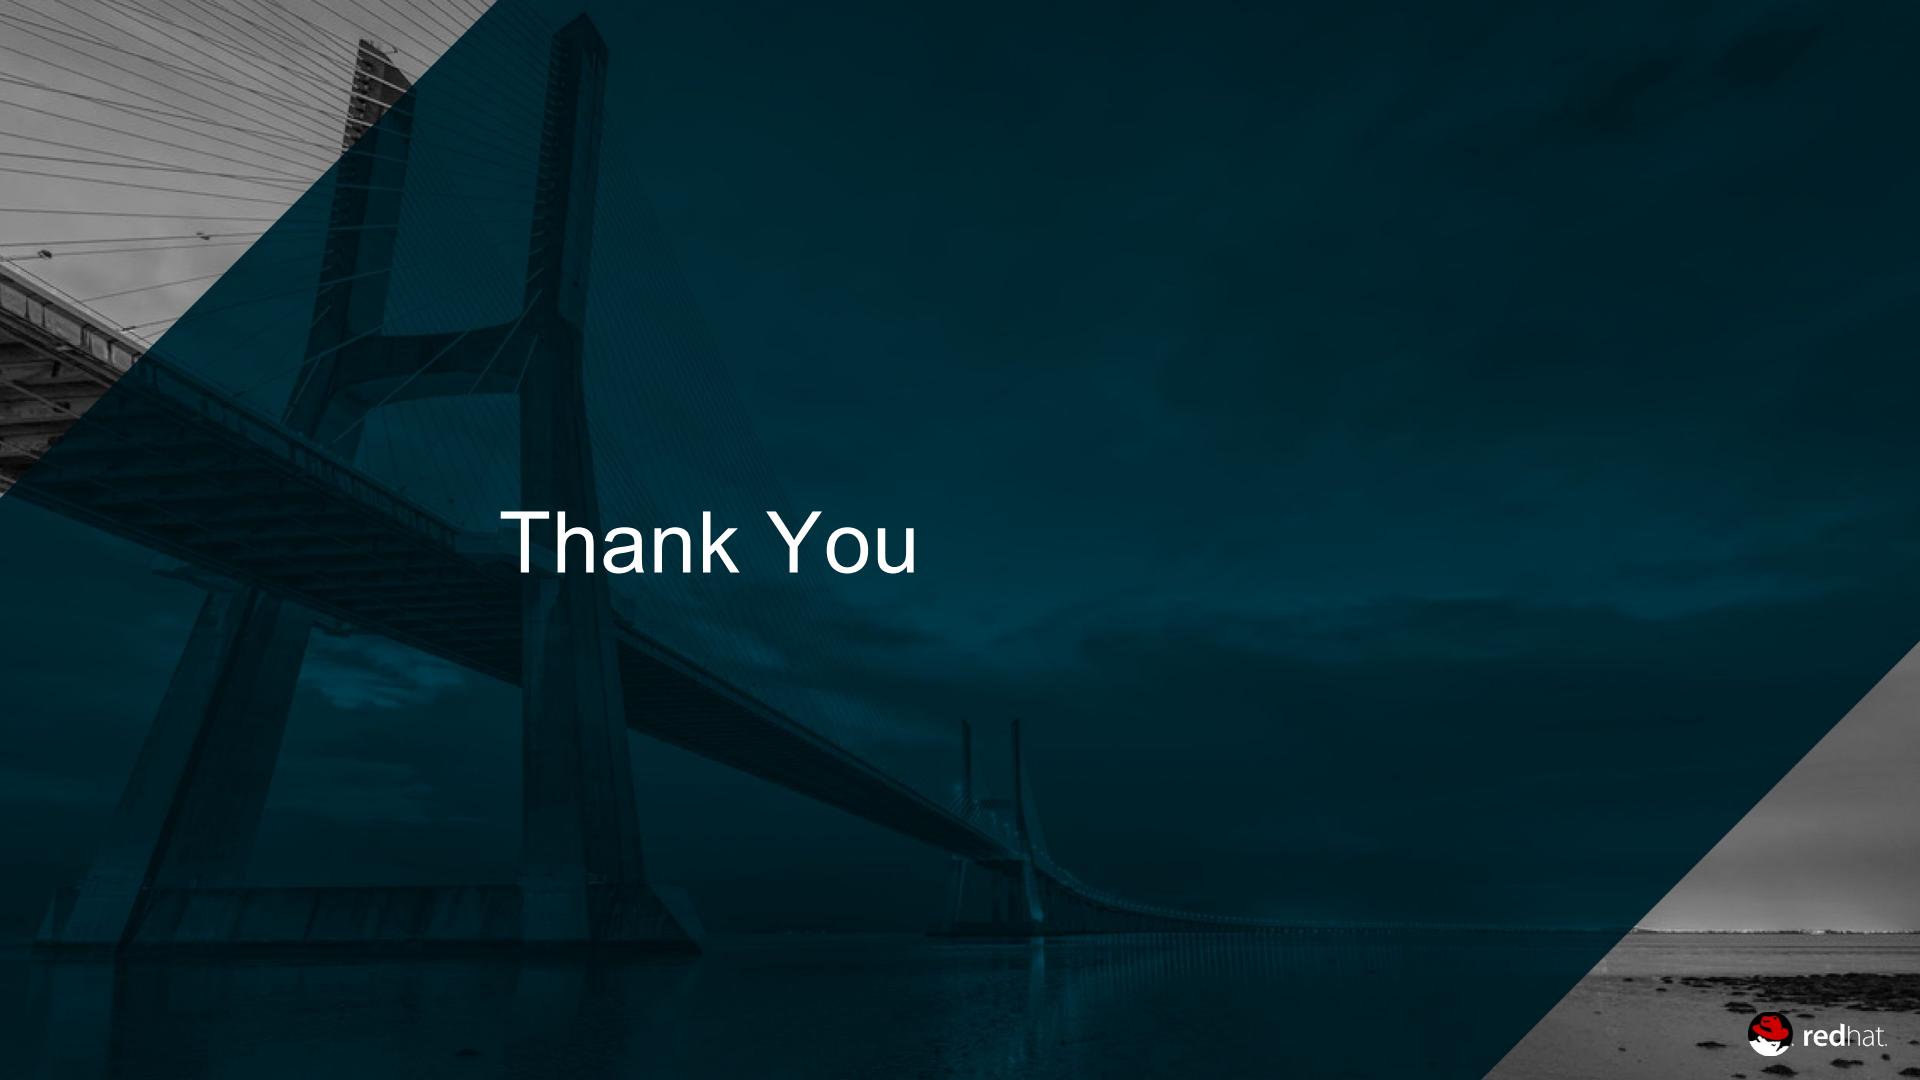
- To create the files necessary for a new chart, you can enter the command:

```
helm create mychart
```



# Demo Placeholder

Let's create our own Charts



# Thank You



# Serverless on Kubernetes

# What is Serverless

---

- Public Cloud providers often provide Serverless capabilities in which you can **deploy functions**, rather than instances or containers
  - Azure Functions
  - AWS Lambda
  - Google Cloud Functions
- With these products, you **don't need to manage the underlying infrastructure**
- The functions are also **not “always-on”** unlike containers and instances, which can greatly reduce the cost of serverless if the function doesn't need to be executed a lot

# What is Serverless

- Serverless in public cloud can **reduce the complexity, operational costs, and engineering time to get code running**
  - You don't need to manage a Windows/Linux distribution
  - You don't need to build containers
  - You only pay for the time your function is running
  - A developer can “just push” the code and does not worry about many operational aspects
    - Although “cold-starts”, the time for a function to start after it has not been invoked for some time, can be an operational issue that needs to be taken care of

# What is Serverless

---

- This is an example of a AWS Lambda Function:

```
exports.handler = function(event, context) {  
    context.succeed("Hello, World!");  
};
```

- You'd still need to setup when the code is being executed
- For example in AWS you would use the API Gateway, to setup a URL that will invoke this function when visited

# Serverless in Kubernetes

- Rather than using containers to start applications on Kubernetes, you can also use **Functions**
- Currently, the **most popular projects** enabling functions are:
  - OpenFaas
  - Kubeless
  - Fission
  - OpenWhisk
- You can install and use any of the projects to let developers launch functions on your Kubernetes cluster
- As an **administrator**, you'll **still need to manage the underlying infrastructure**, but from a **developer** standpoint, he/she will be able to **quickly and easily deploy functions on Kubernetes**

# Kubeless

# Kubeless

---

- Kubeless is a Kubernetes-native framework (source: <https://github.com/kubeless/kubeless/>)
  - It leverages the Kubernetes resources to provide auto-scaling, API routing, monitoring, etc
- It uses **Custom Resource Definitions** to be able to create functions
- It's **open source** and **non-affiliated** to any commercial organization
- It has a UI available for developers to deploy functions

# Kubeless

---

- With kubeless you deploy a function in your preferred language
- Currently, the **following runtimes are supported:**
  - Python
  - NodeJS
  - Ruby
  - PHP
  - .NET
  - Golang
  - Others

# Kubeless

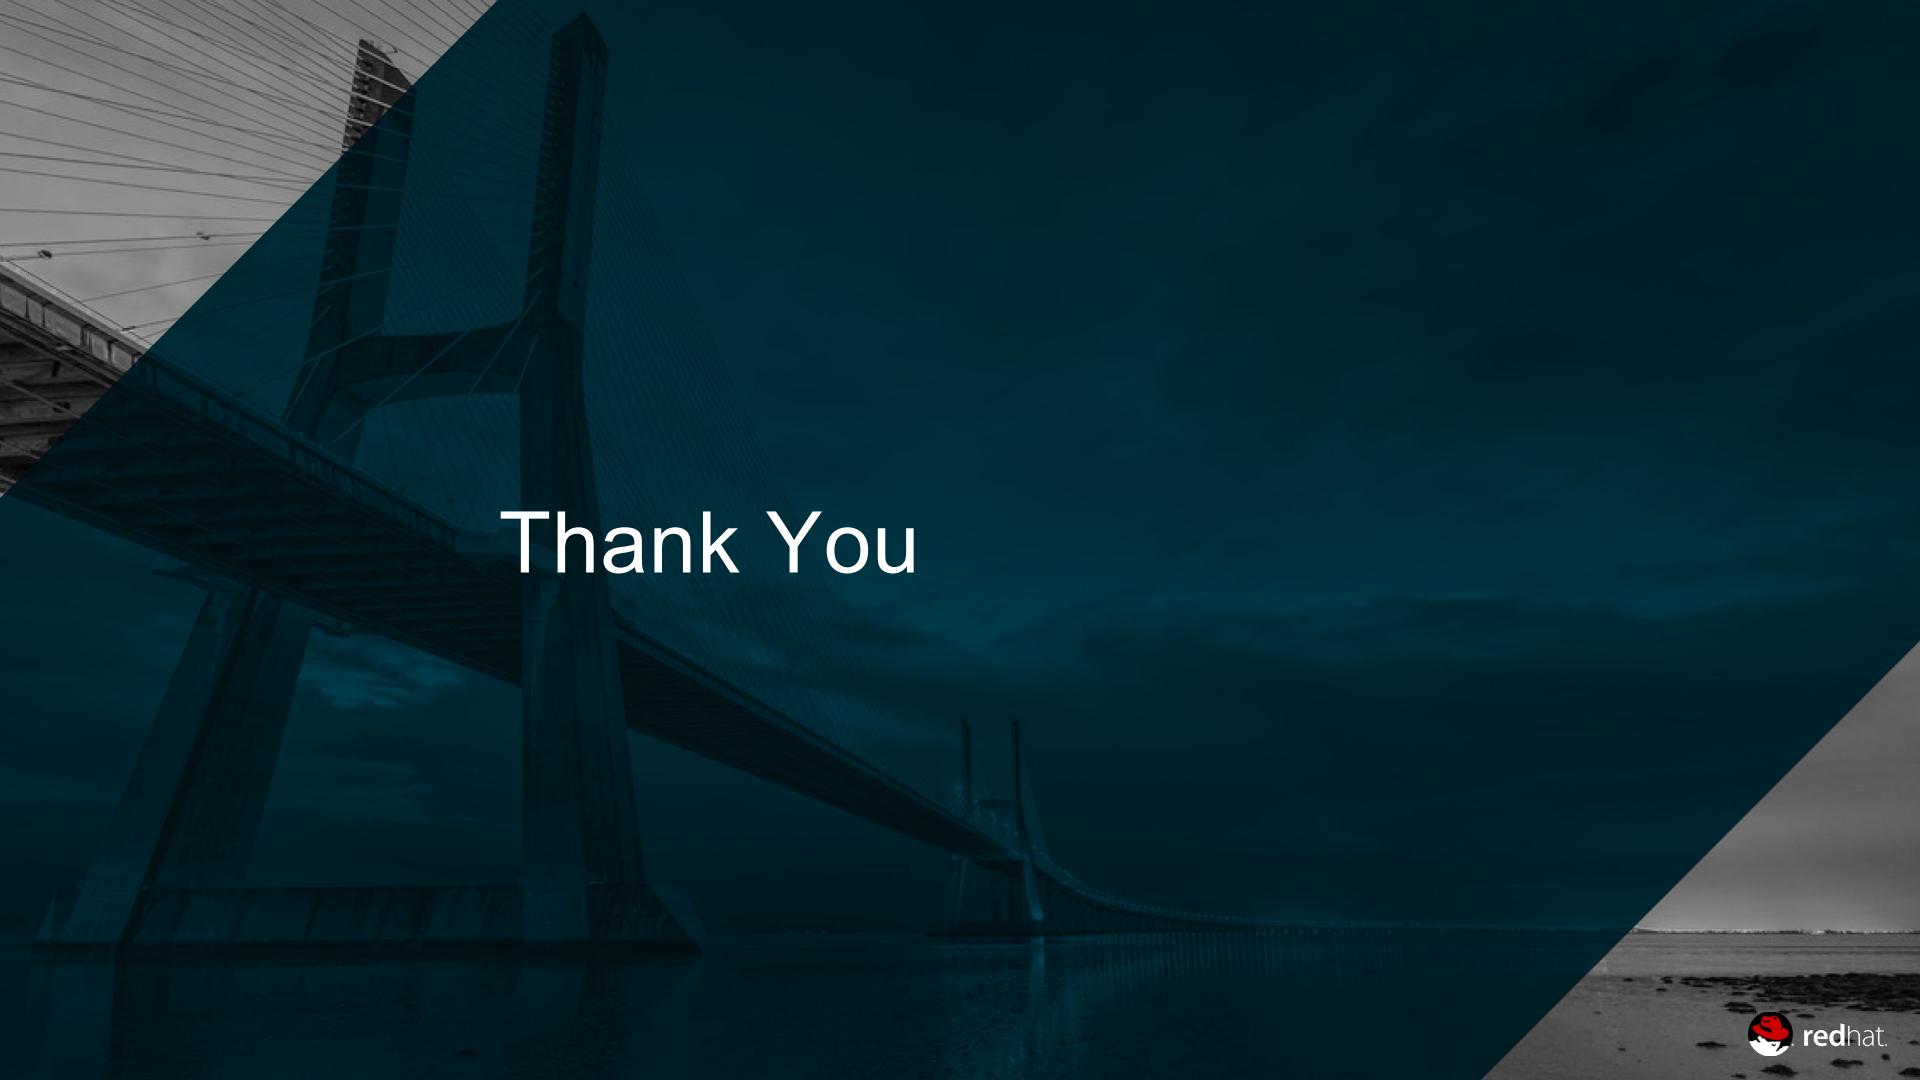
- Once you deployed your function, you'll need to determine how it'll be **triggered**
- Currently, the **following function are supported**:
  - HTTP functions**
    - HTTP functions gets executed when an HTTP endpoint is triggered
    - You write a function and return the text/HTML that needs to be displayed in the browser
  - Scheduled function

# Kubeless

- Once you deployed your function, you need to determine how it'll be **triggered**
- Currently, the **following function are supported:**
  - PubSub (Kafka or NATS)
    - Triggers a function when data is available in Kafka / NATS
  - AWS Kinesis
    - Triggers based on data in AWS Kinesis (similar to Kafka)

# Demo Placeholder

Let's create our own Charts



# Thank You

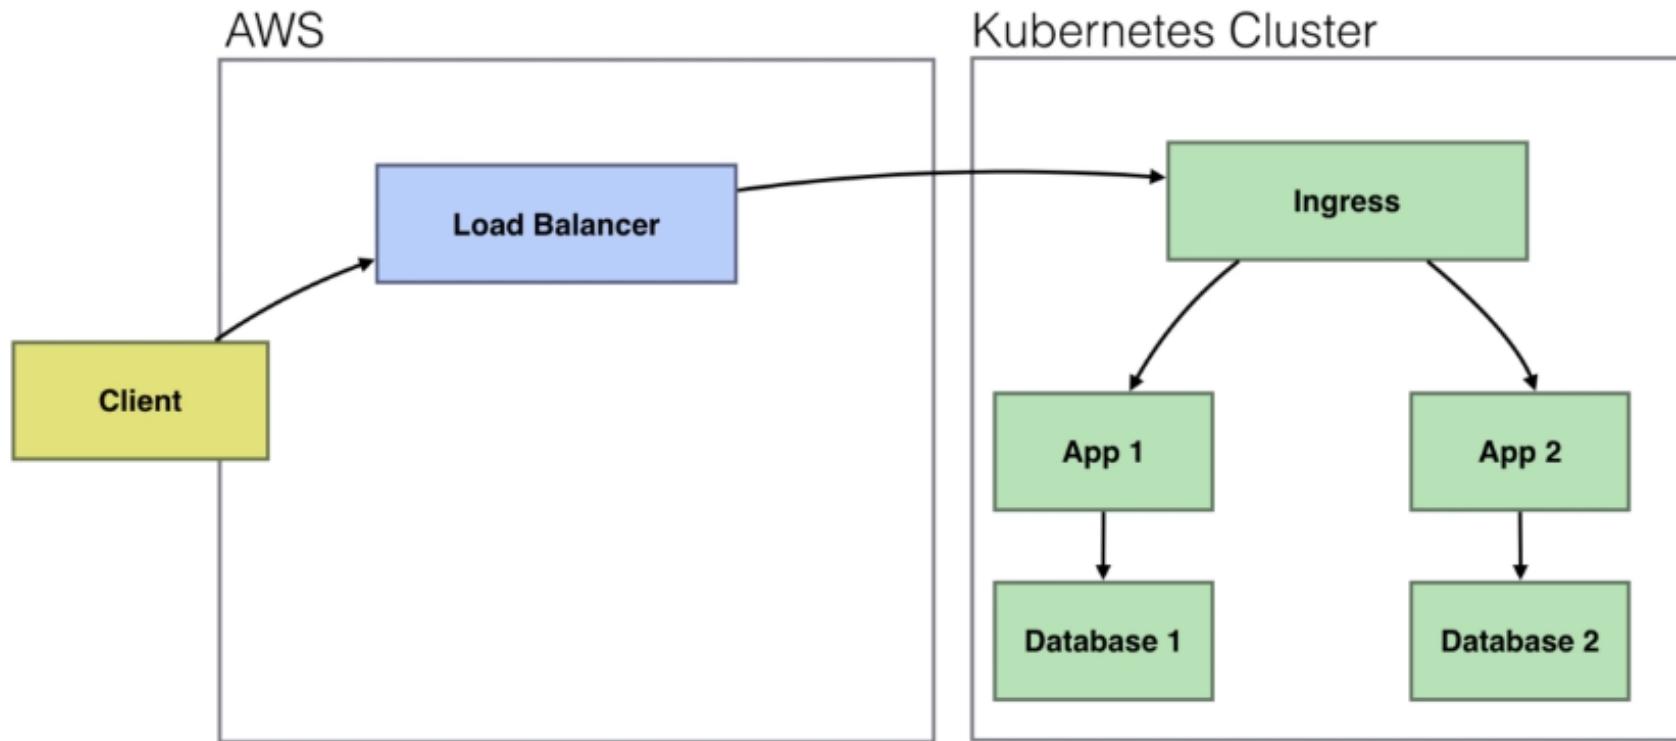
# MicroServices

# Microservices

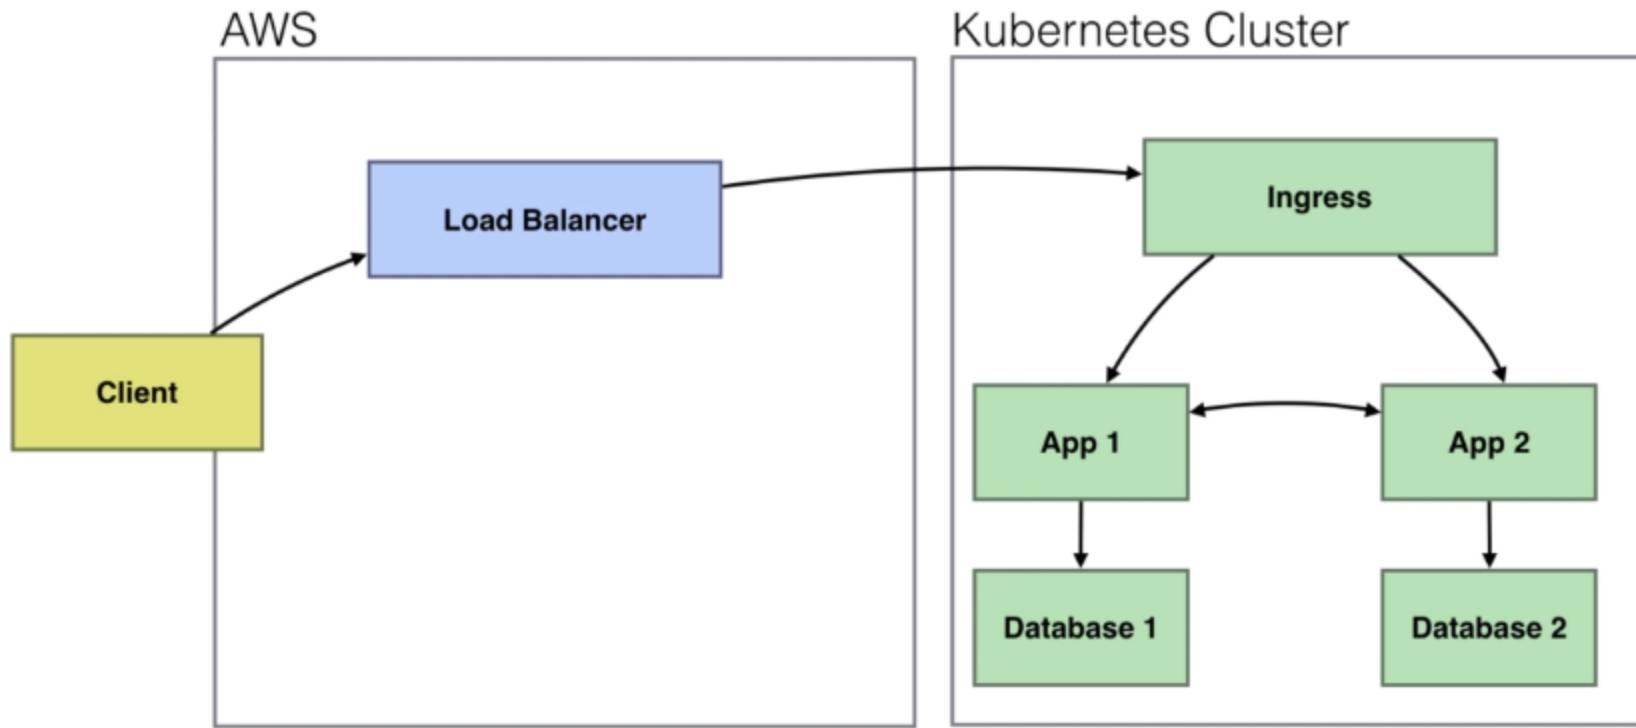
---

- Kubernetes makes it **easy to deploy a lot of diverse applications**
- Those applications can be **monoliths** that don't have anything to do with each other, or **microservices**, small services that make up one application
- The **microservices architecture** is **increasingly popular**
- This approach allows developers to split up the application in **multiple independent parts**
- Having to manage microservices can put an **operational strain** on the engineering team

# Microservices - Monoliths

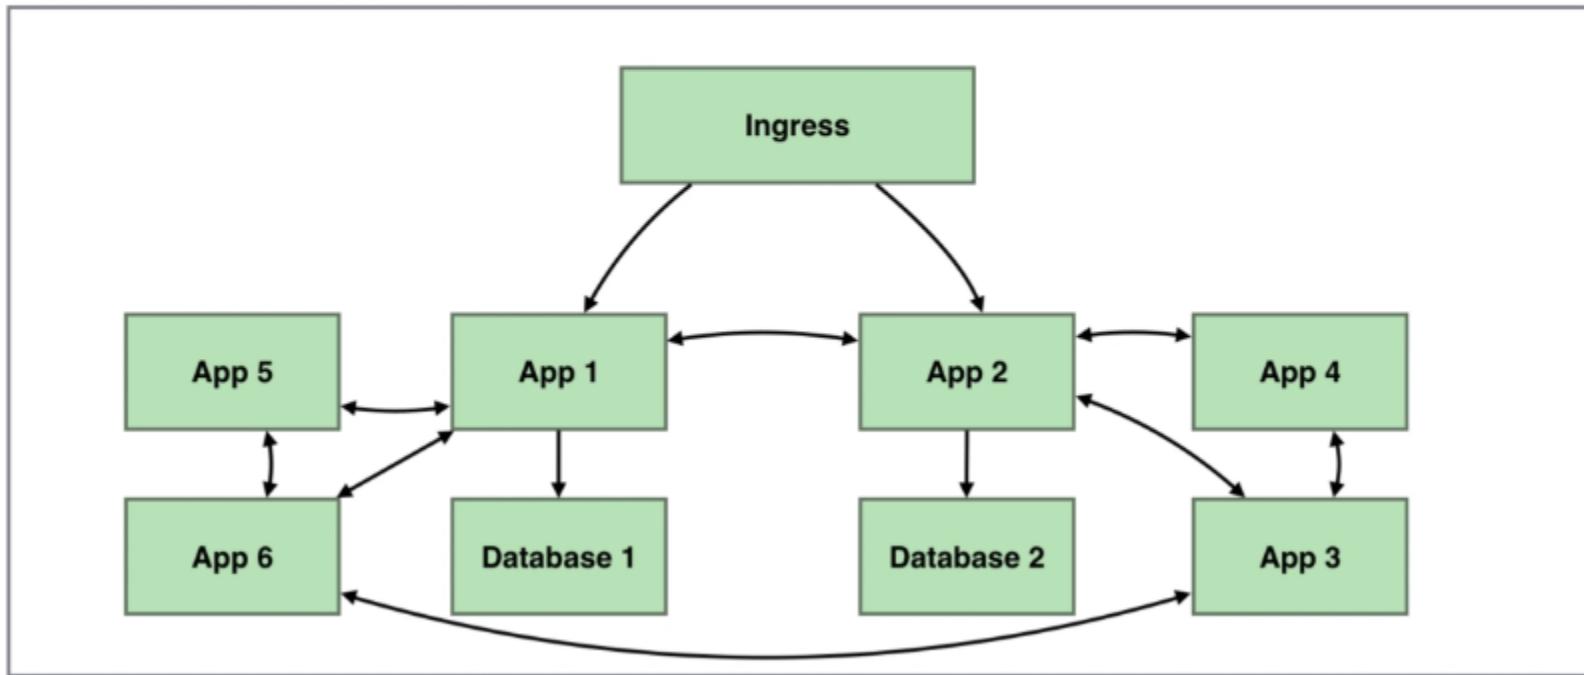


# Microservices



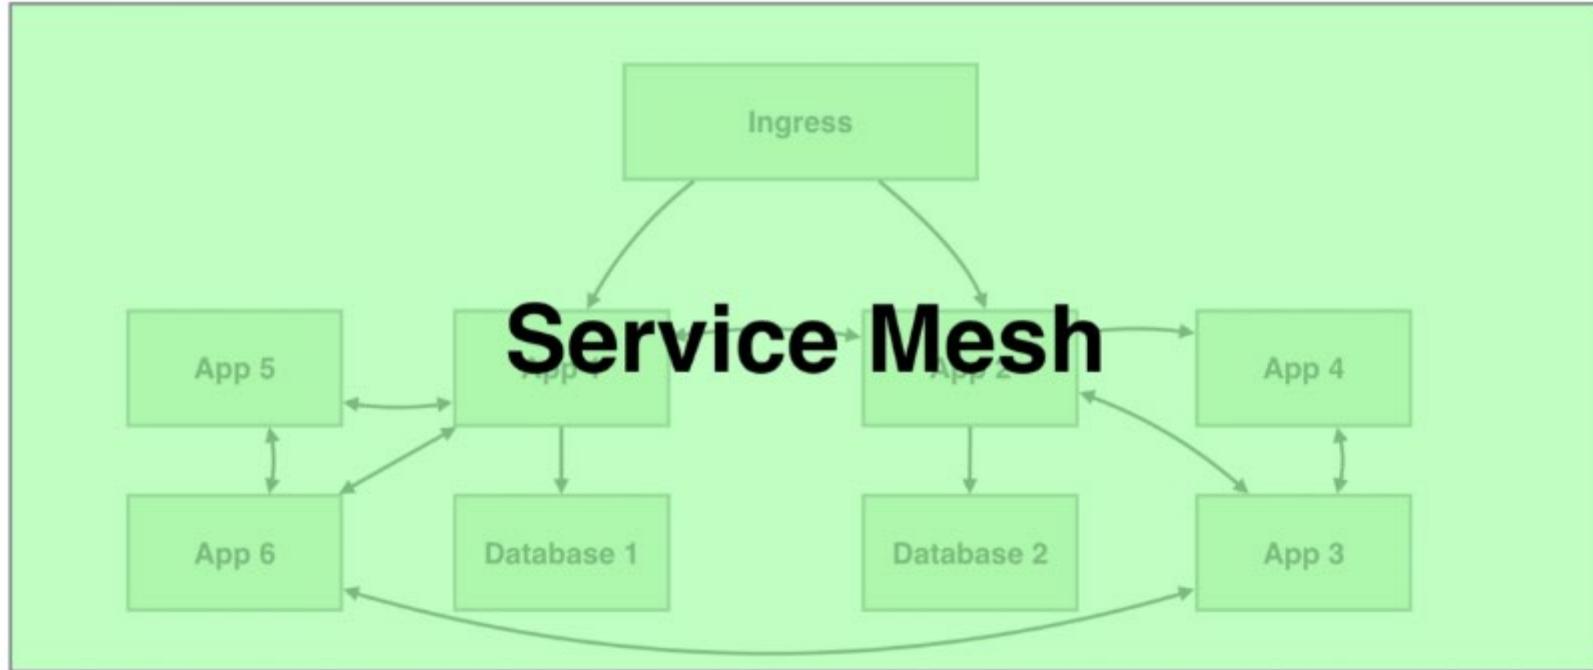
# Microservices

Kubernetes Cluster



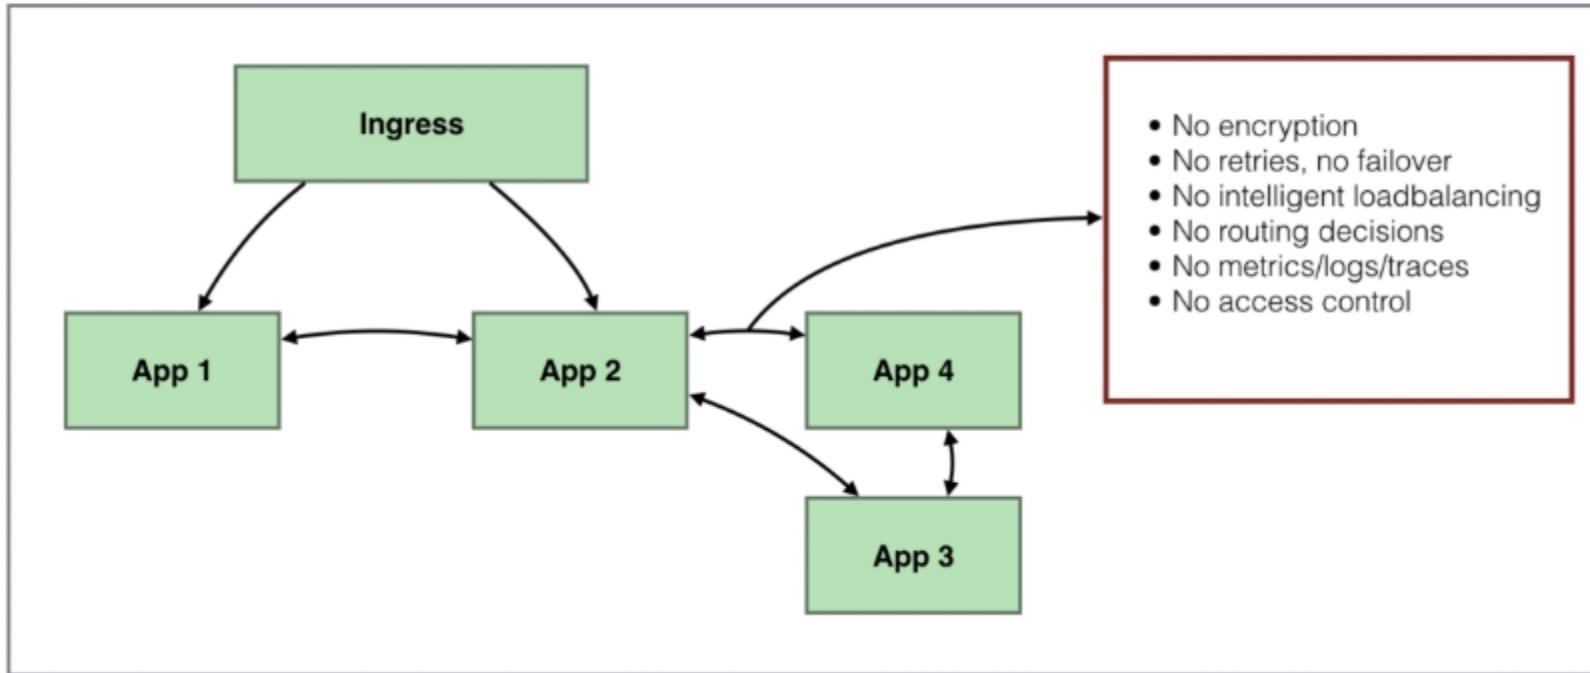
# Microservices

Kubernetes Cluster



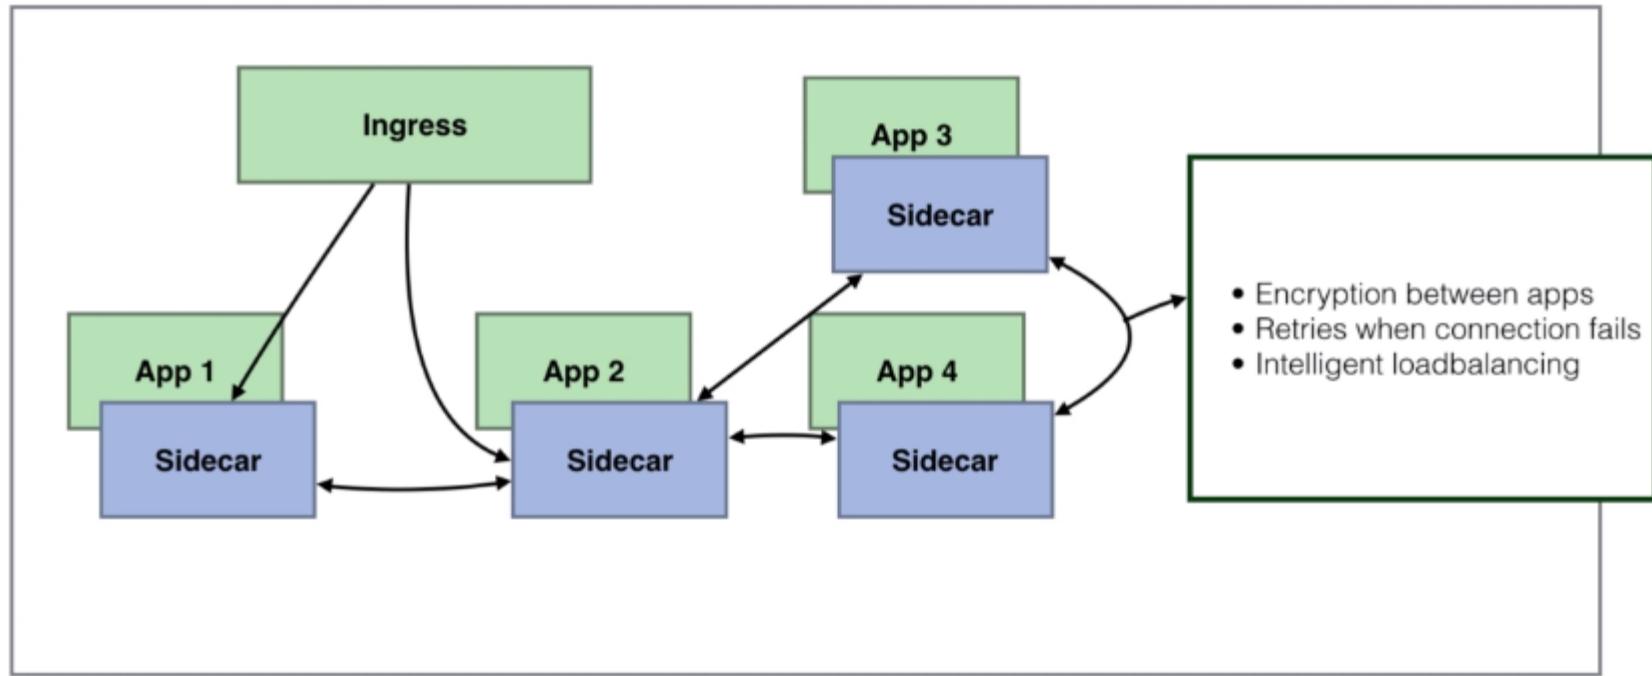
# Microservices

## Kubernetes Cluster



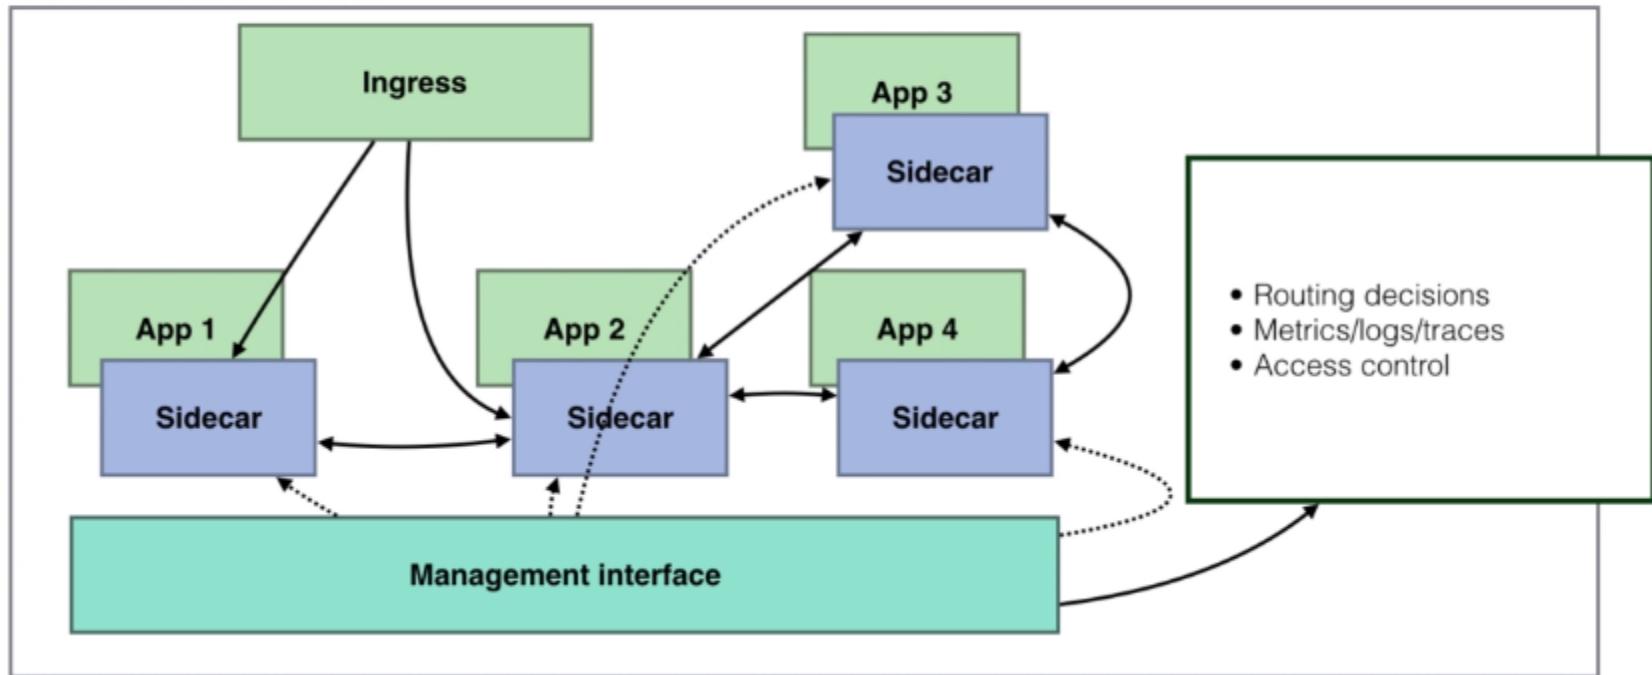
# Microservices

Kubernetes Cluster



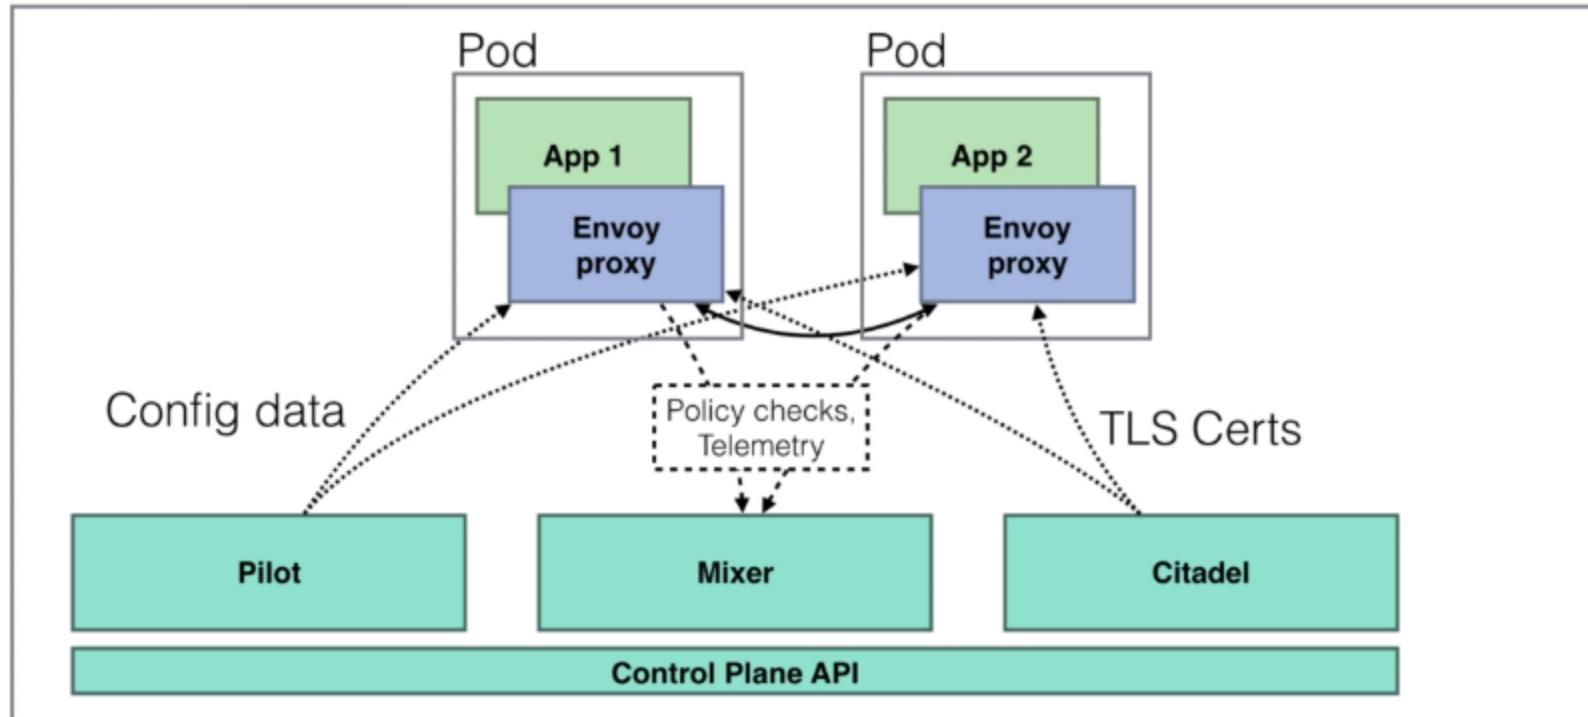
# Microservices

## Kubernetes Cluster



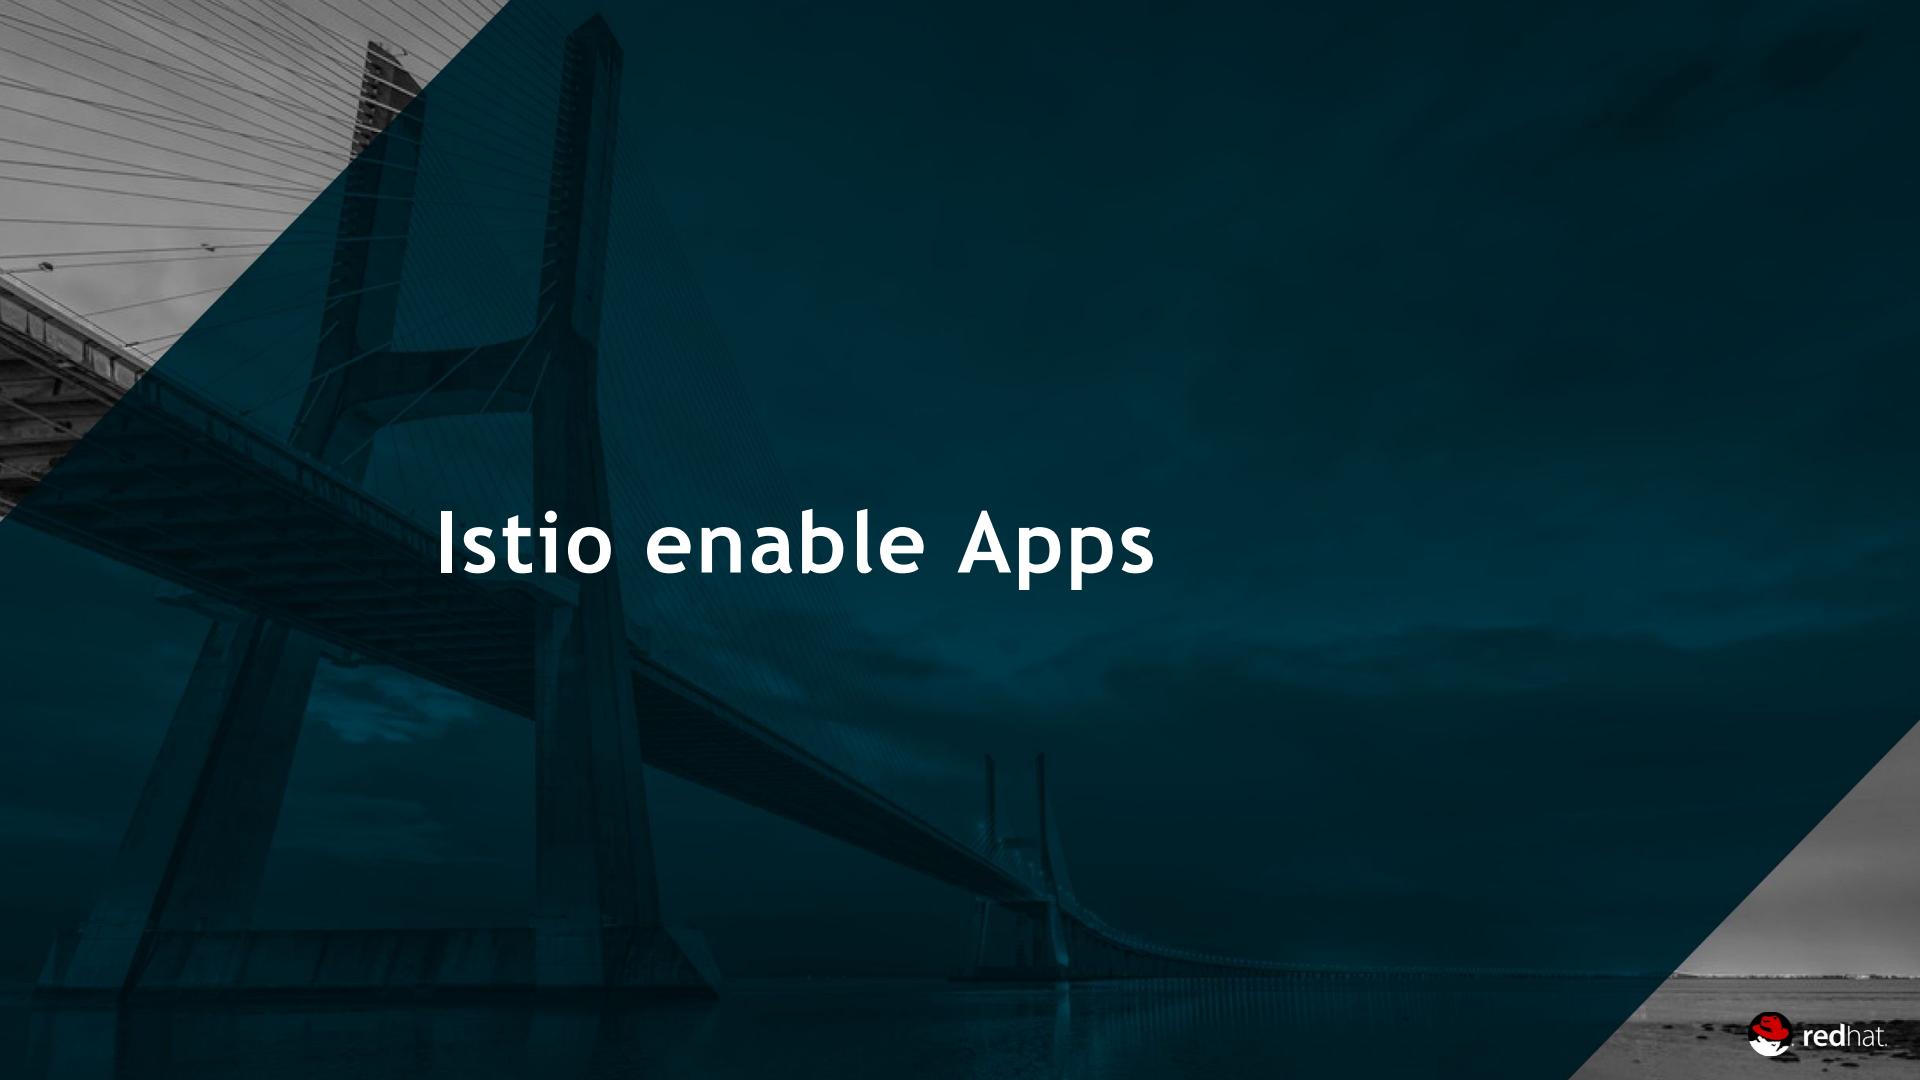
# Istio

Kubernetes Cluster



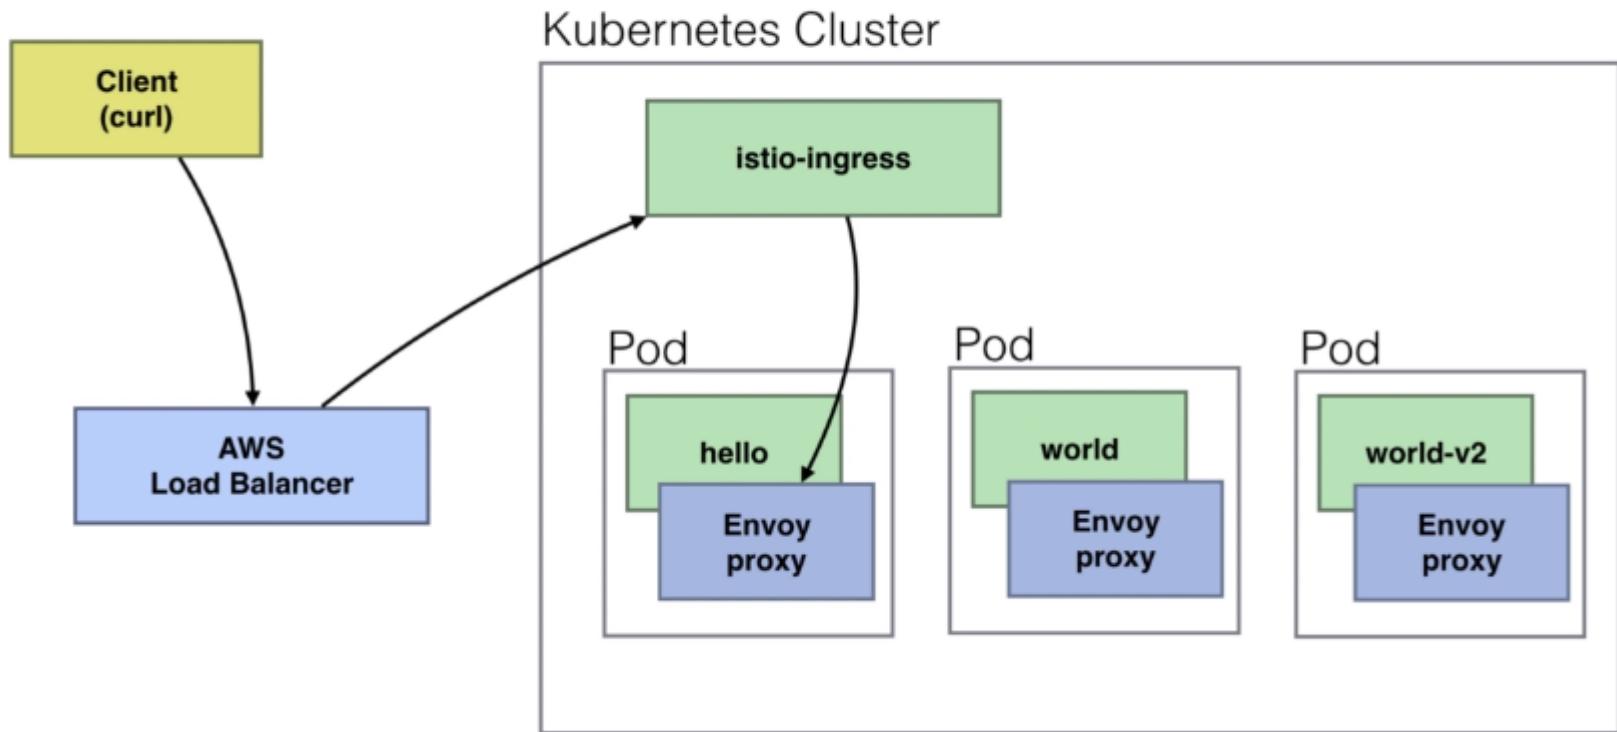
# Demo Placeholder

Istio Installation

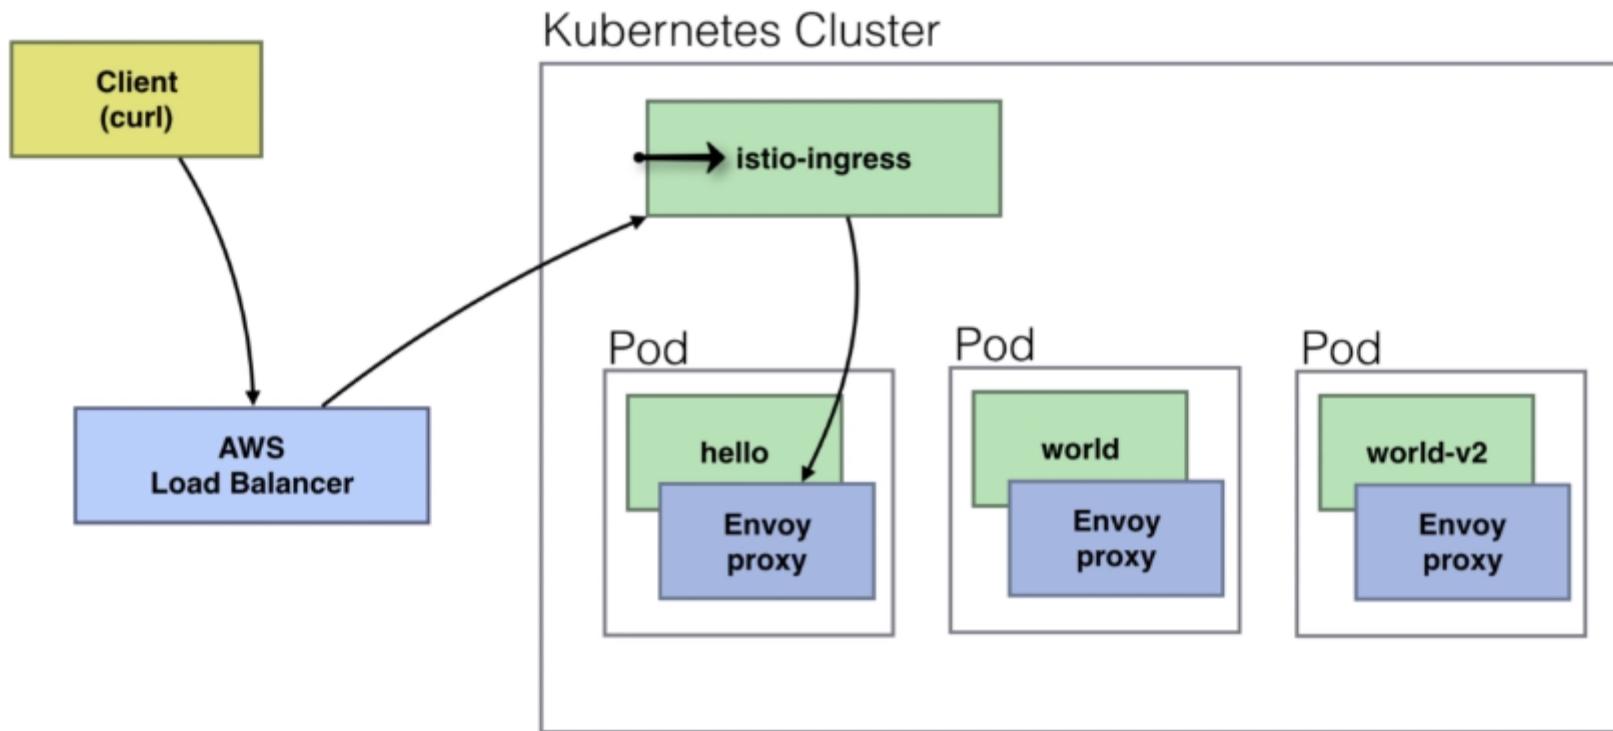


# Istio enable Apps

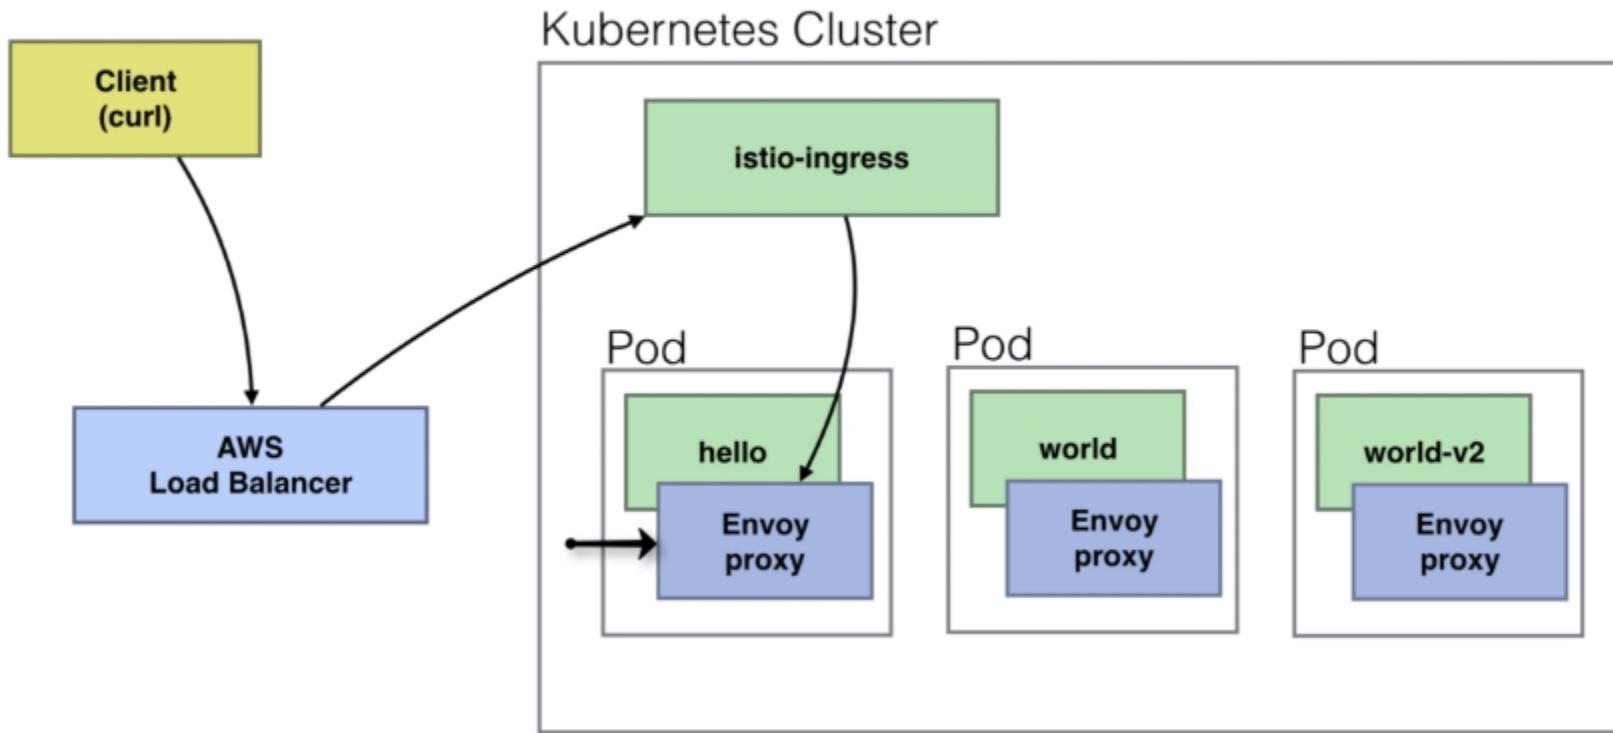
# hello world app



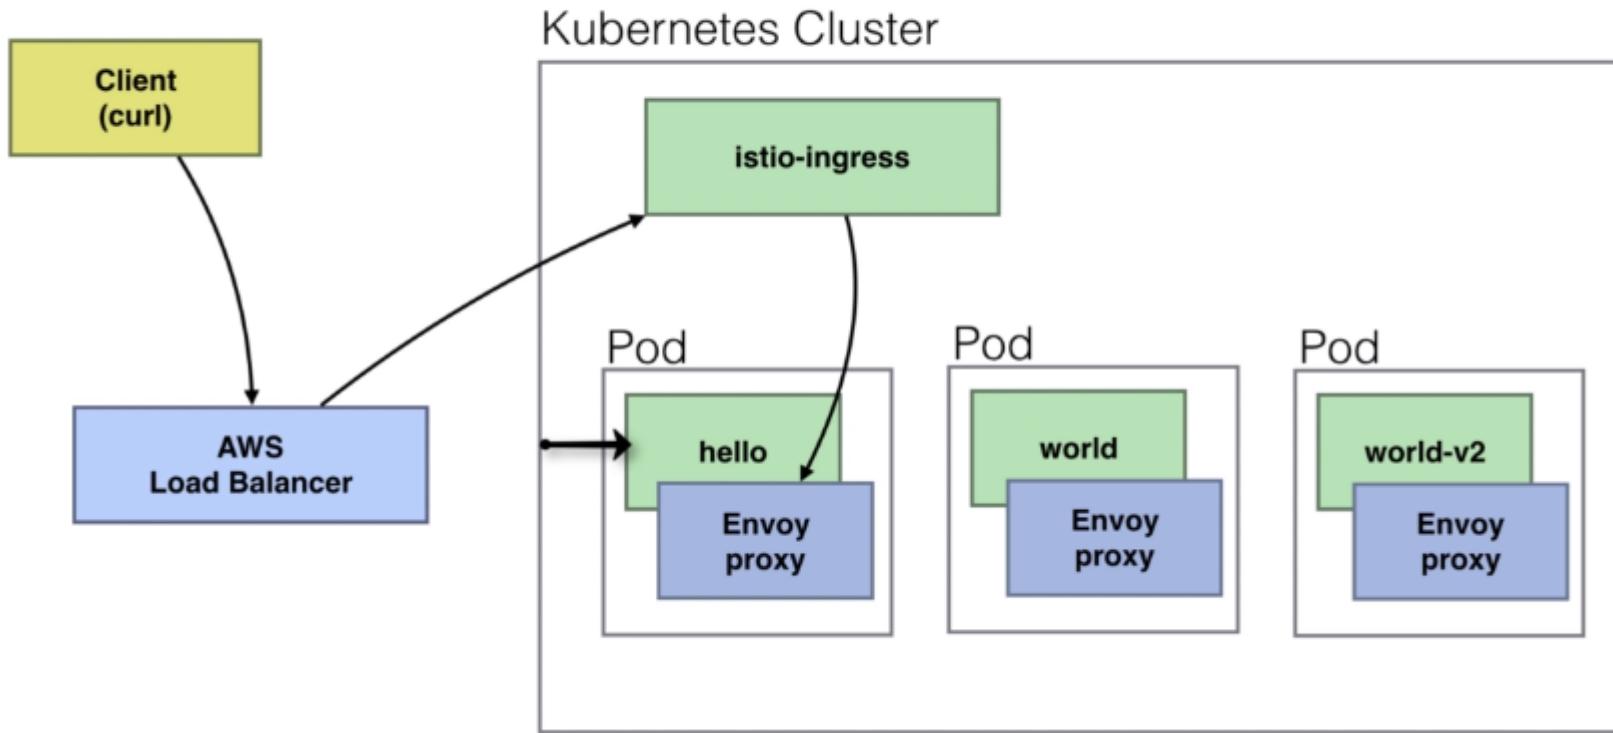
# hello world app



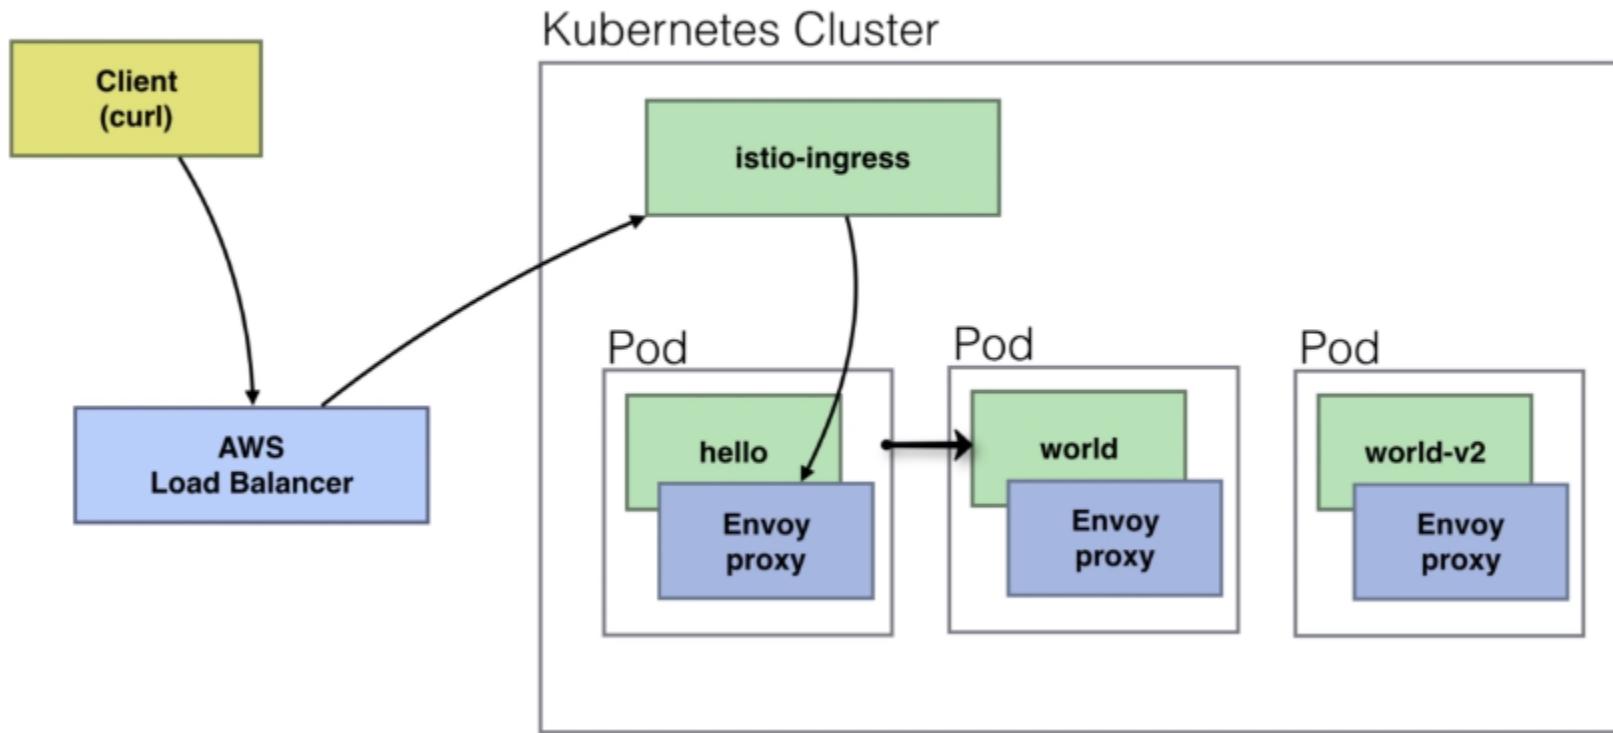
# hello world app



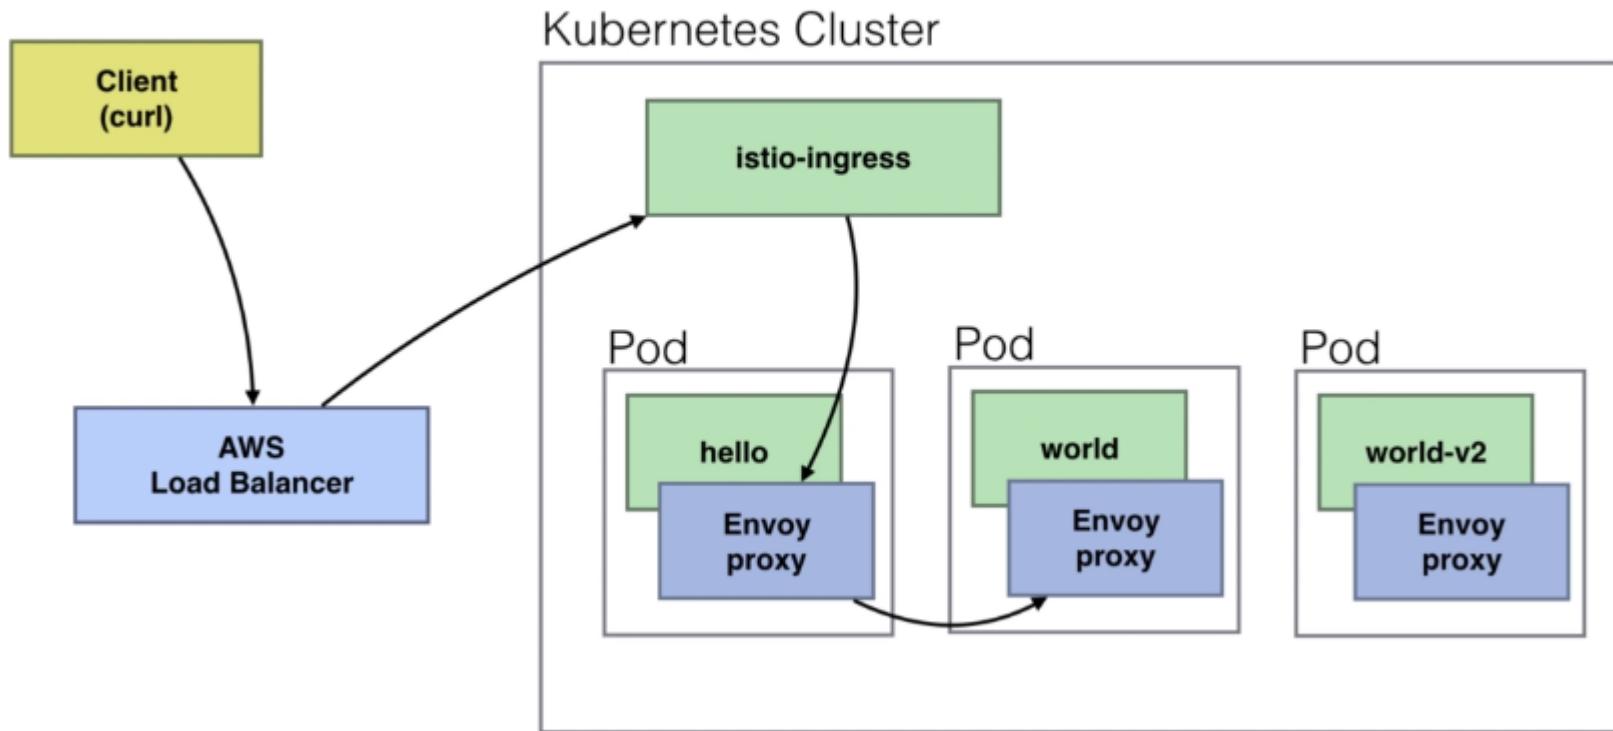
# hello world app



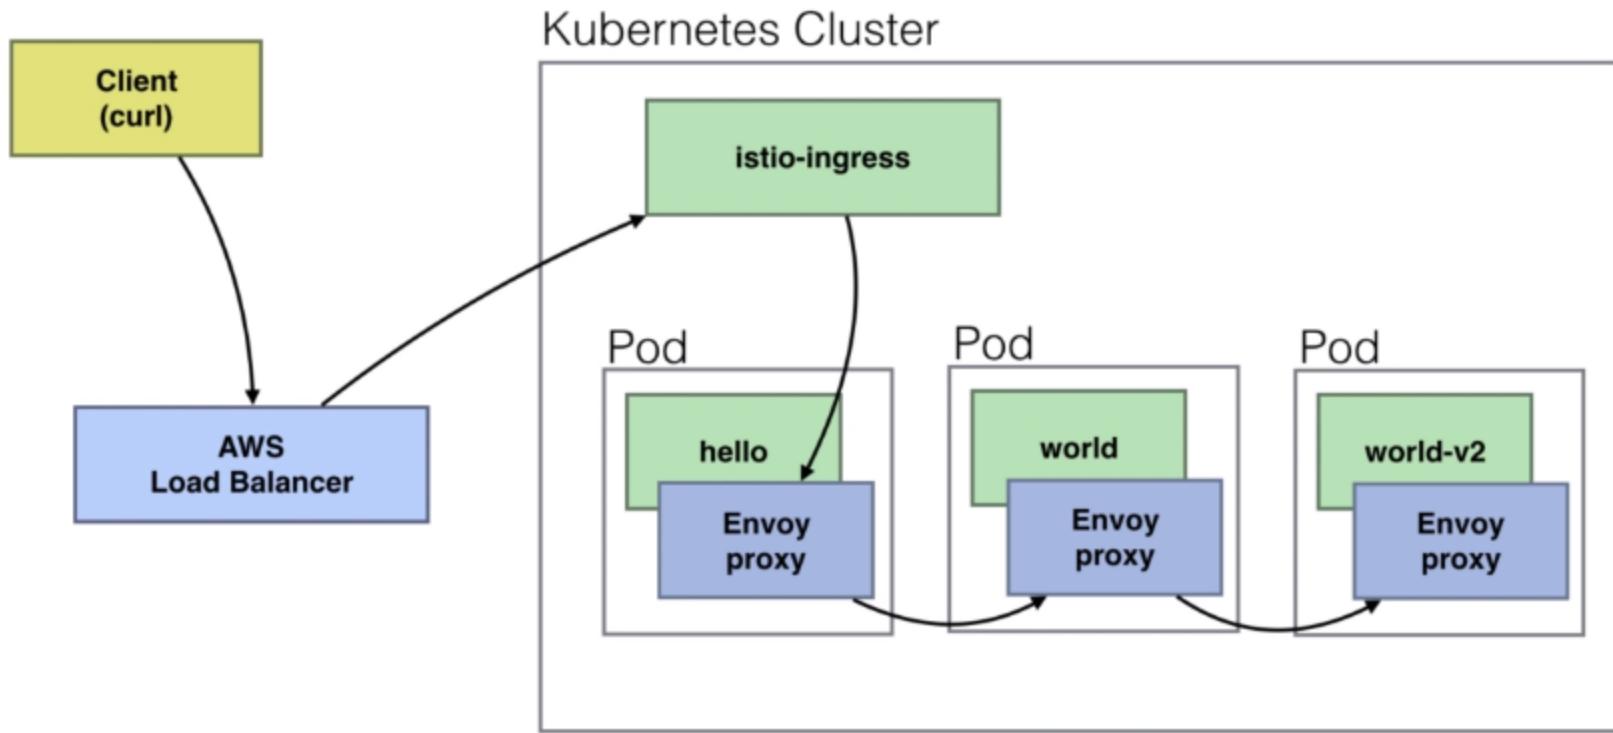
# hello world app



# hello world app



# hello world app



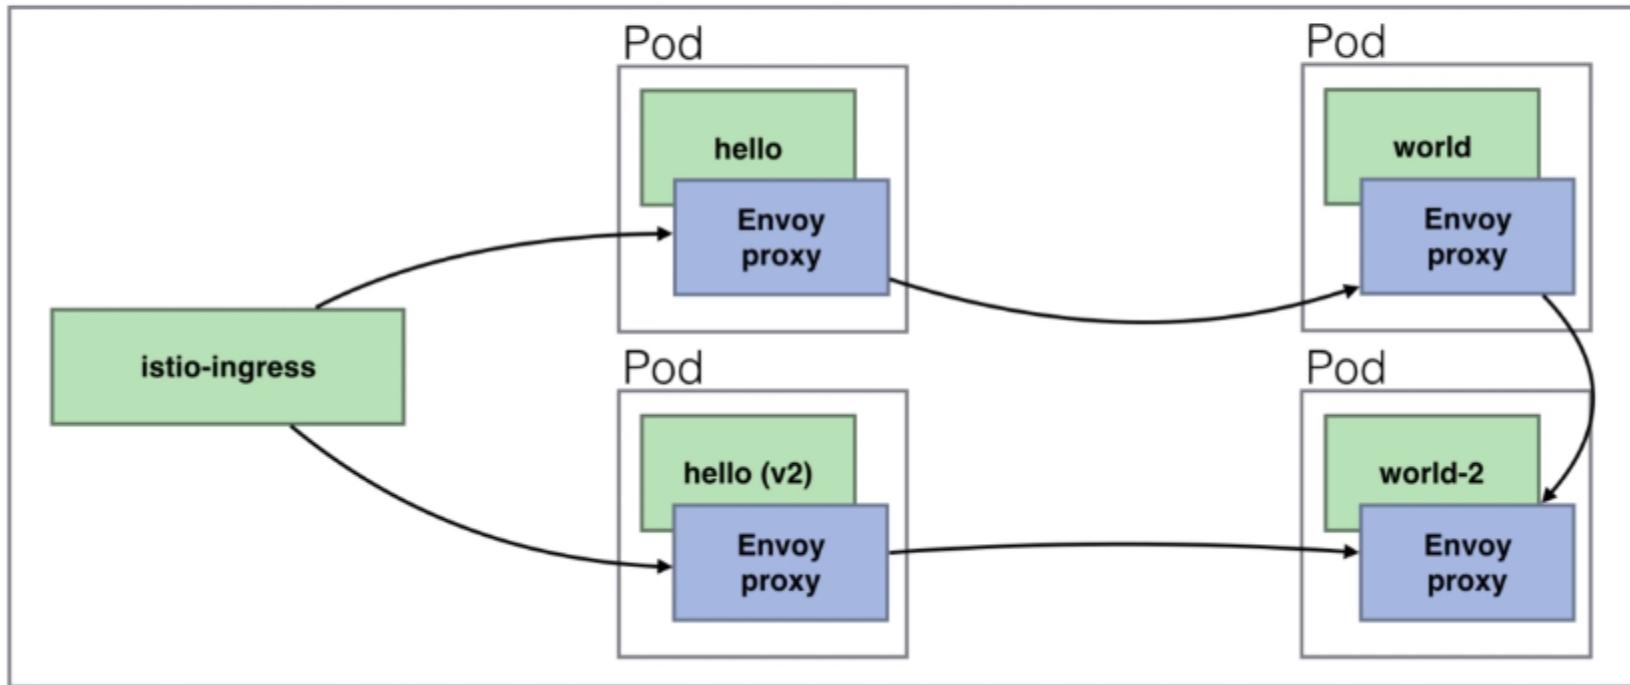
# Demo Placeholder

Setup HelloWord App with istio enabled

# Routing

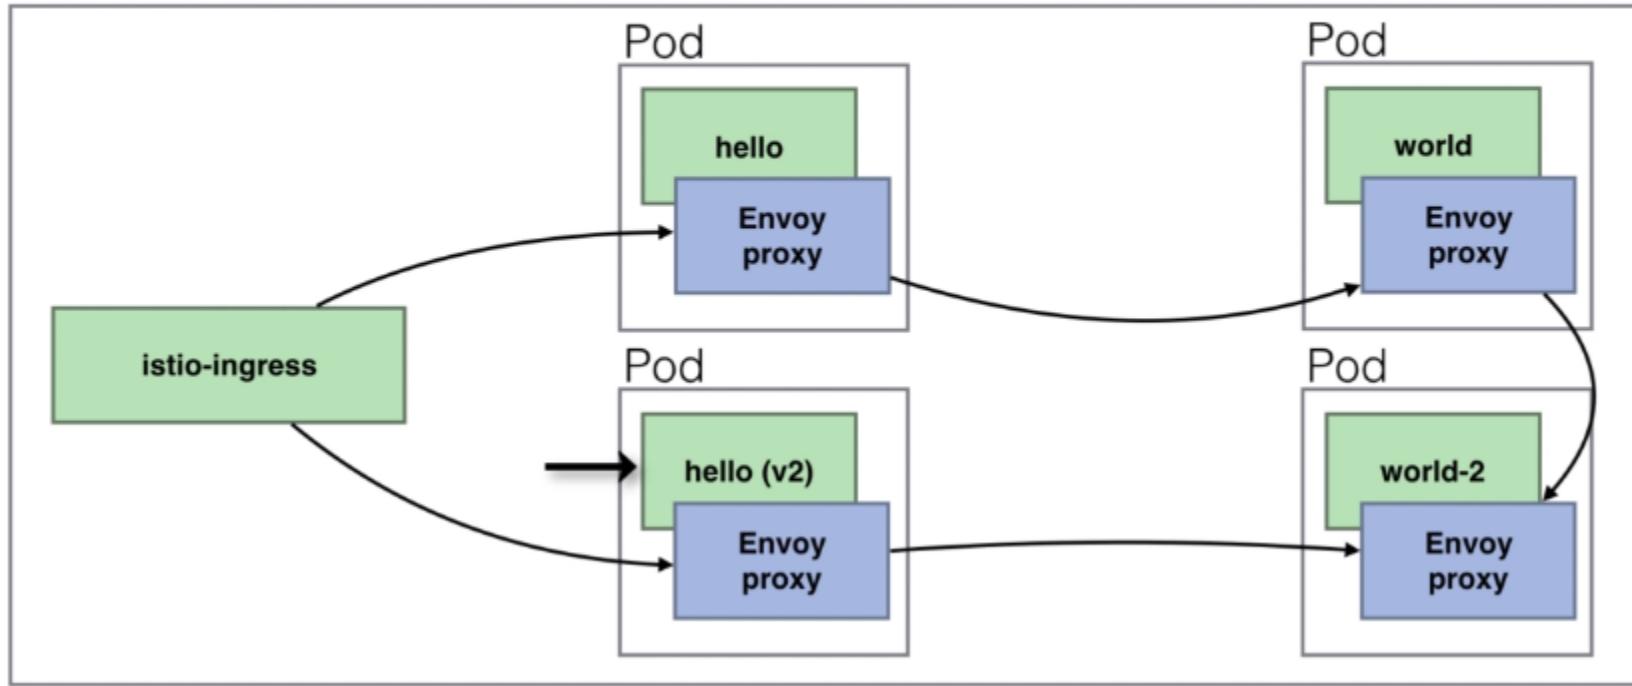
# hello world app - v2

Kubernetes Cluster



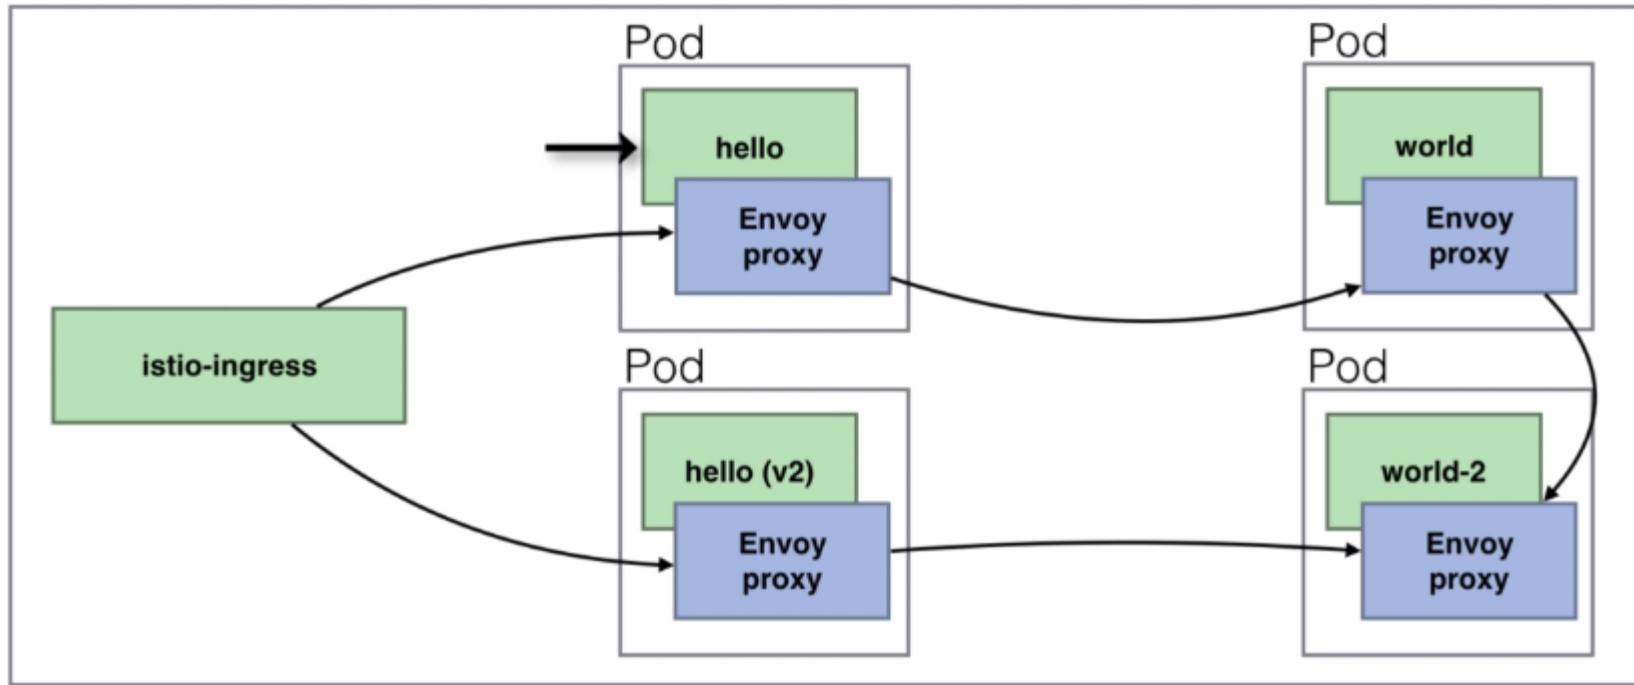
# hello world app - v2

Kubernetes Cluster



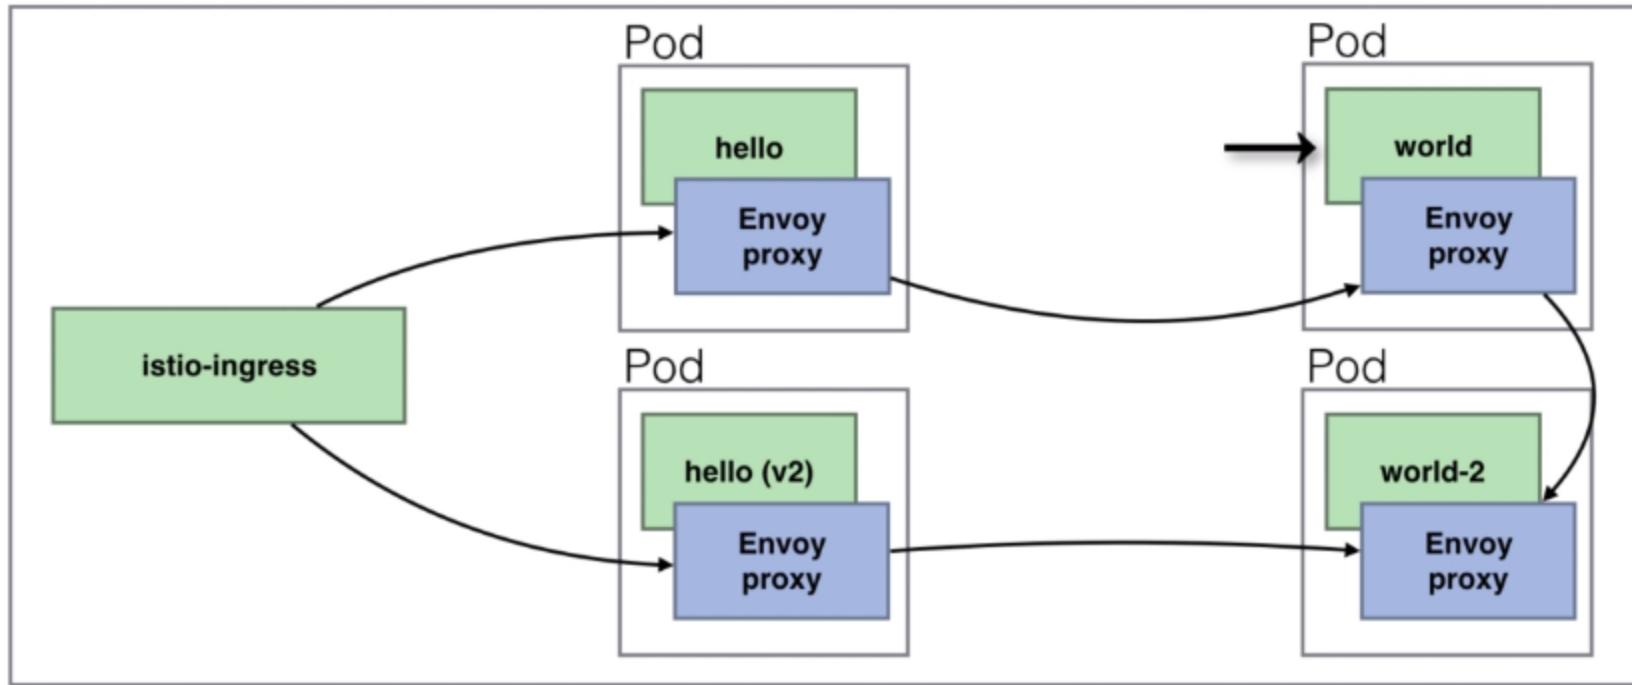
# hello world app - v2

Kubernetes Cluster



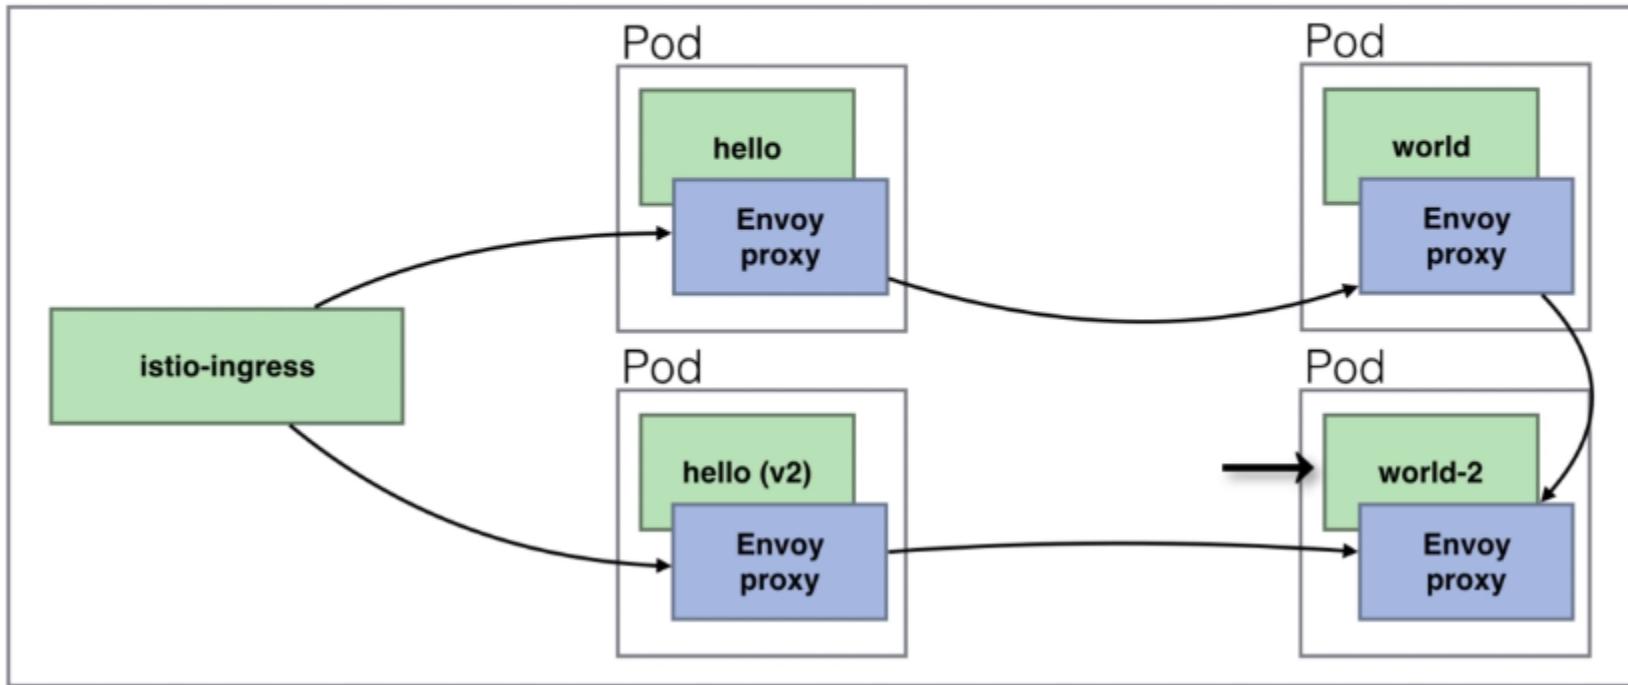
# hello world app - v2

Kubernetes Cluster



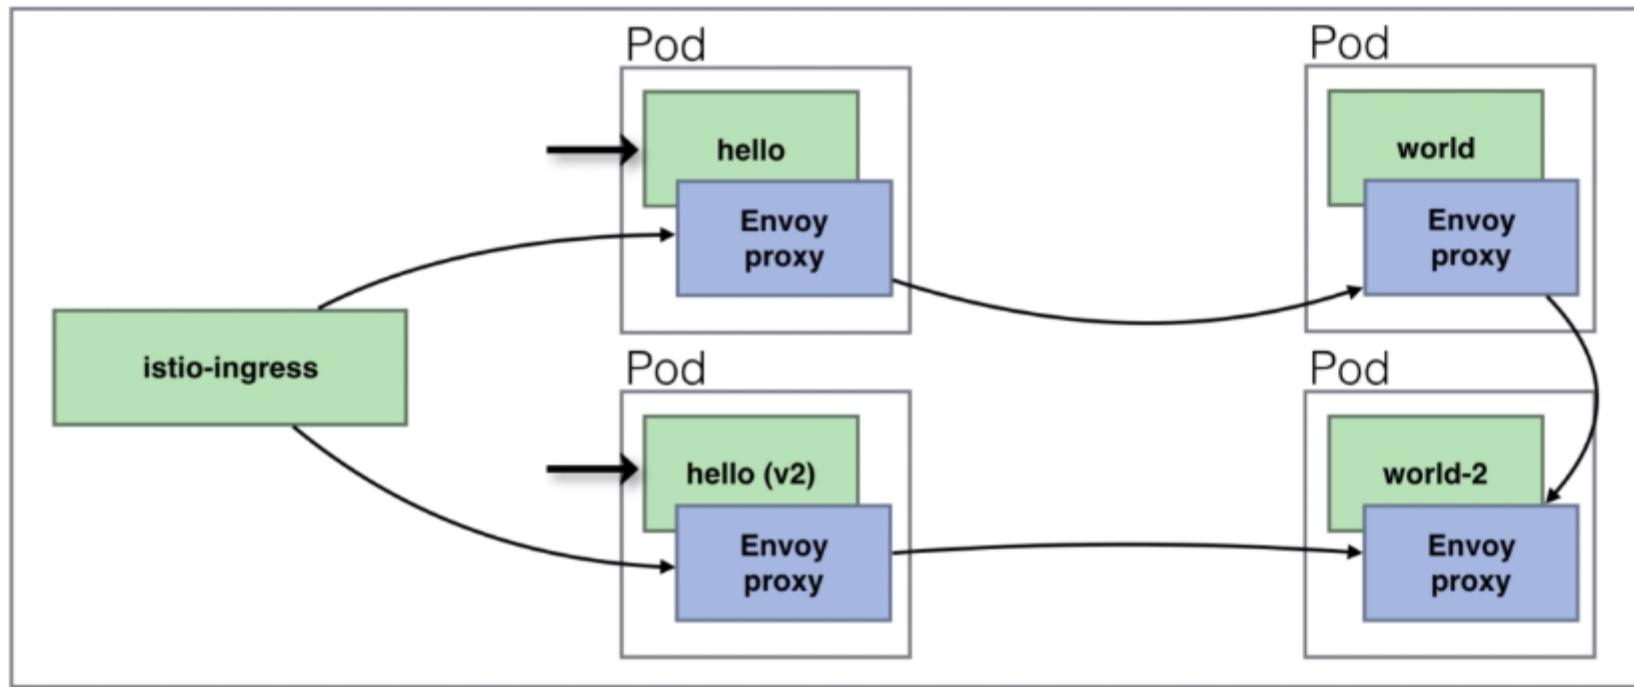
# hello world app - v2

Kubernetes Cluster



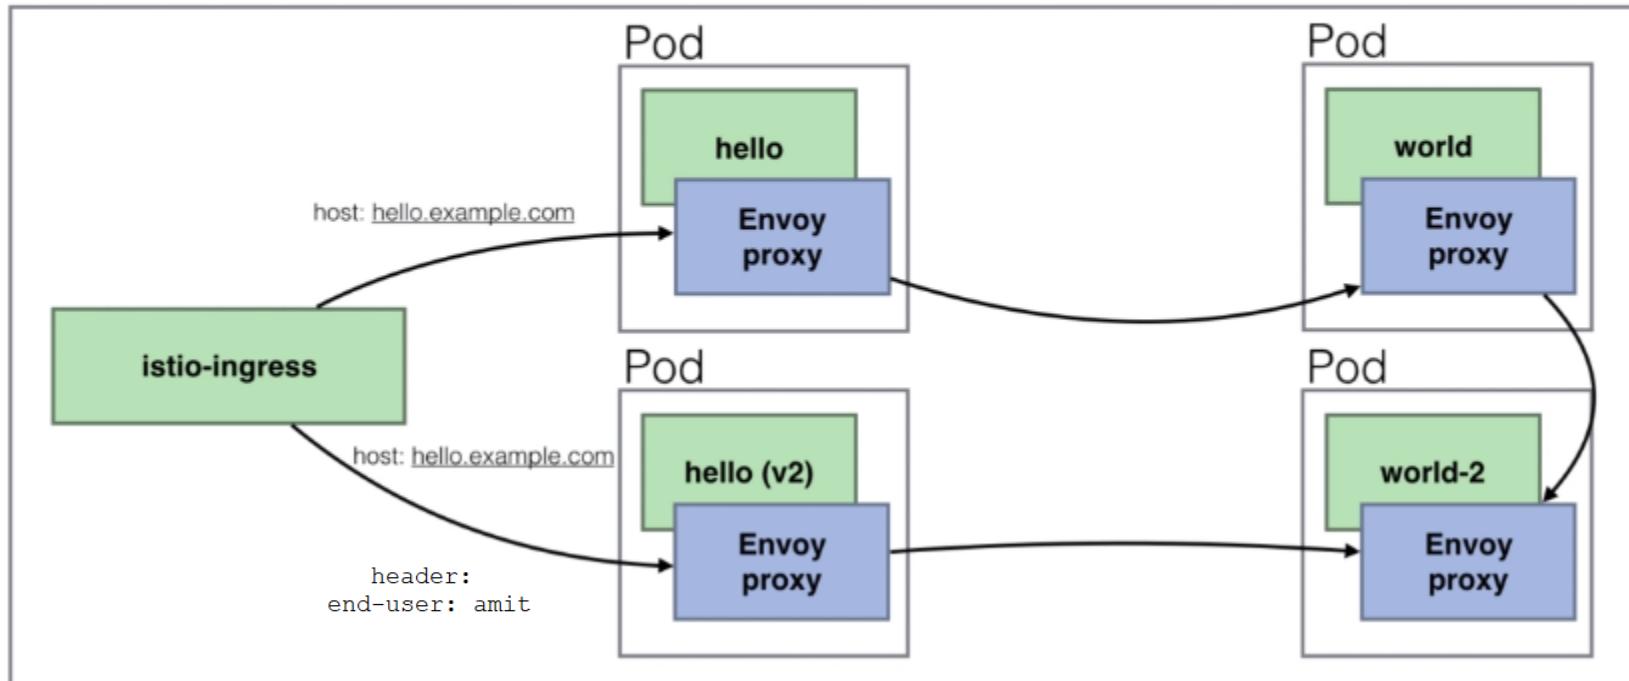
# hello world app - v2

Kubernetes Cluster



# hello world app - v2

Kubernetes Cluster



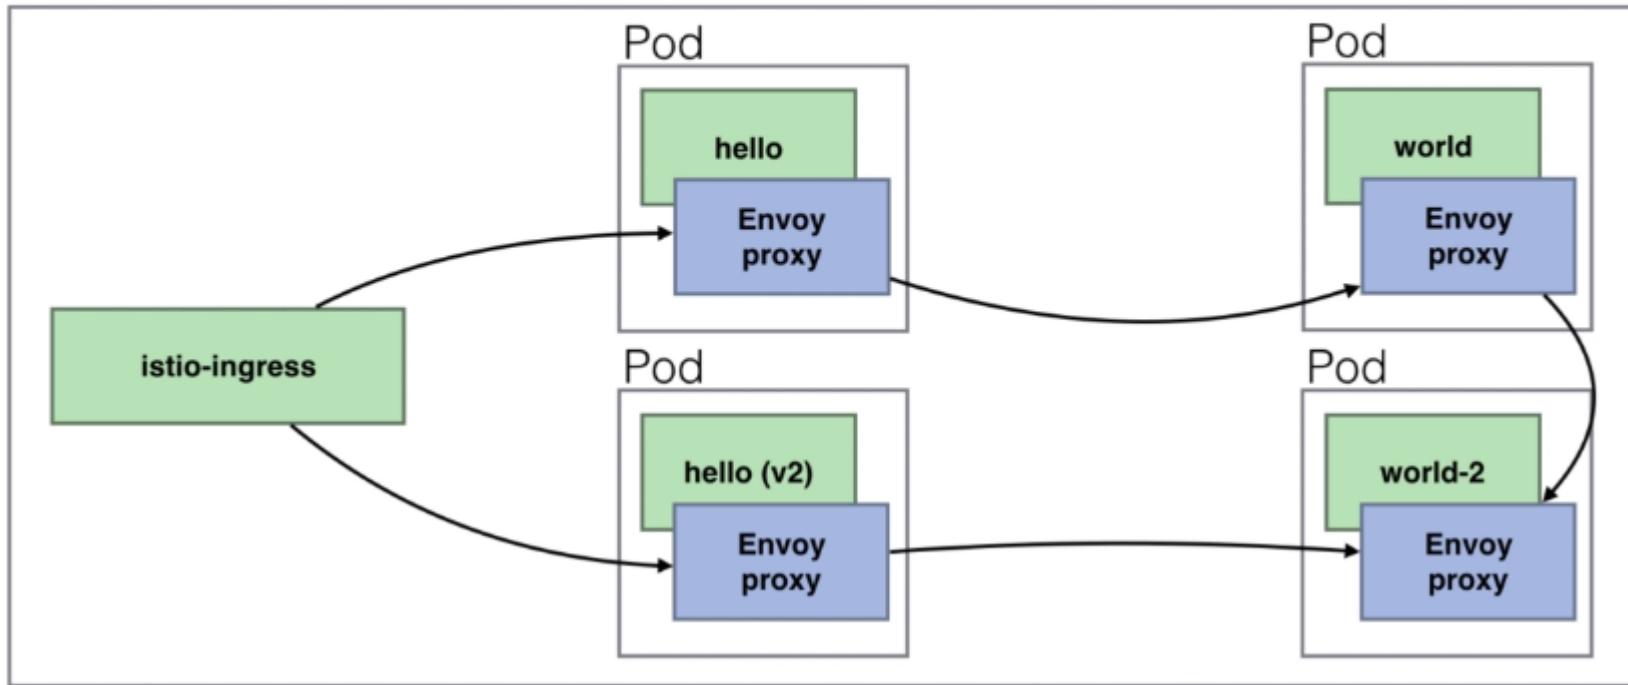
# Demo Placeholder

Hello & Hello-v2 with advance routing

# Canary Deployments

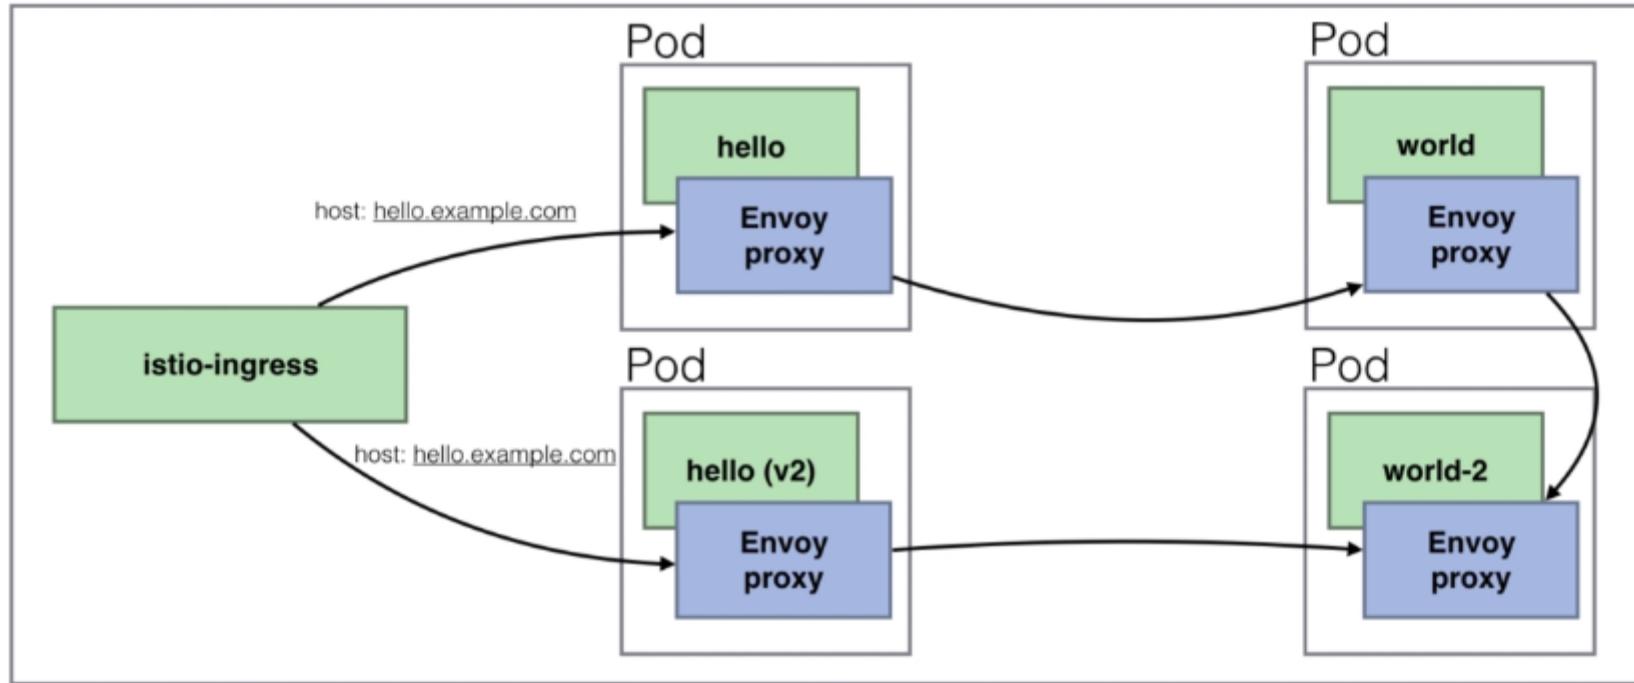
# hello world app - v2

Kubernetes Cluster



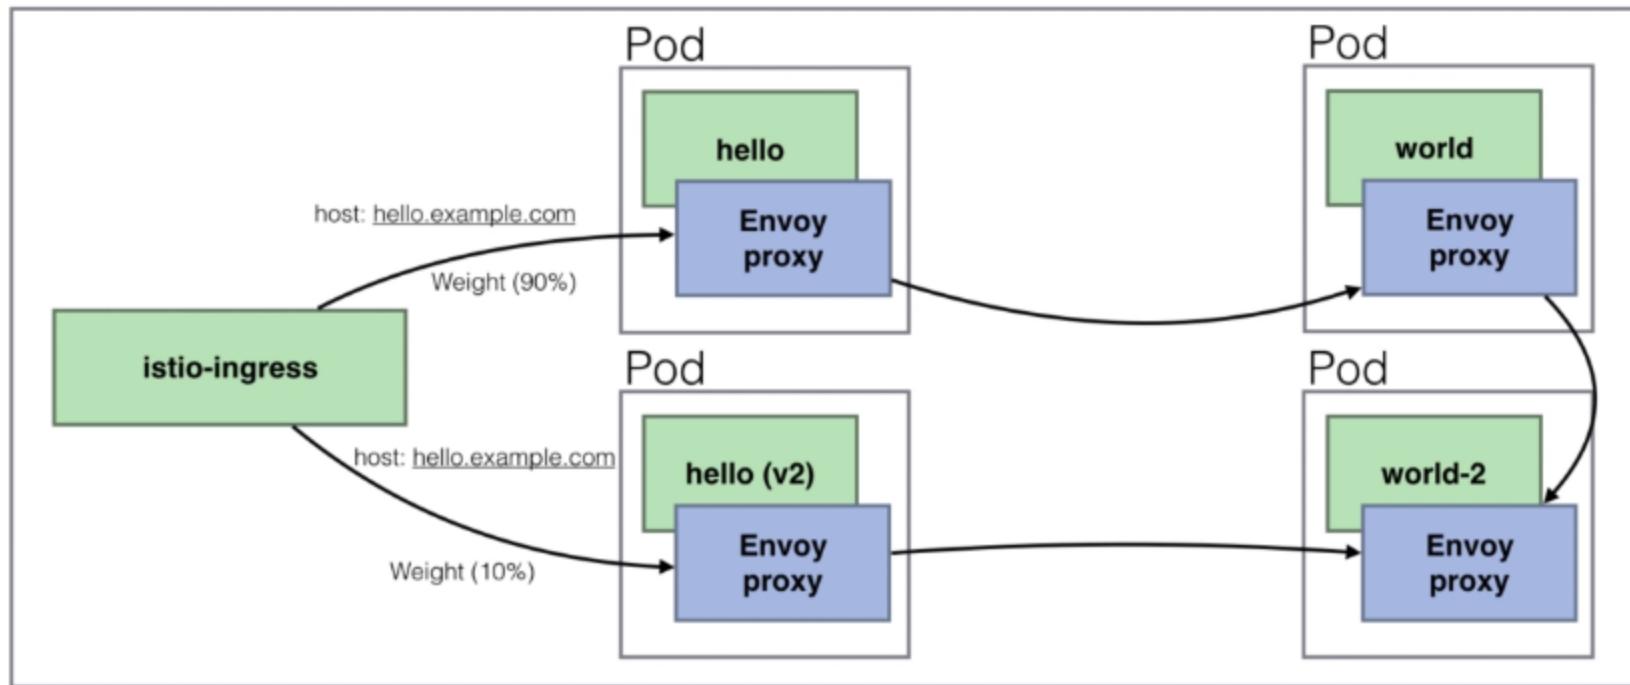
# hello world app - v2

Kubernetes Cluster



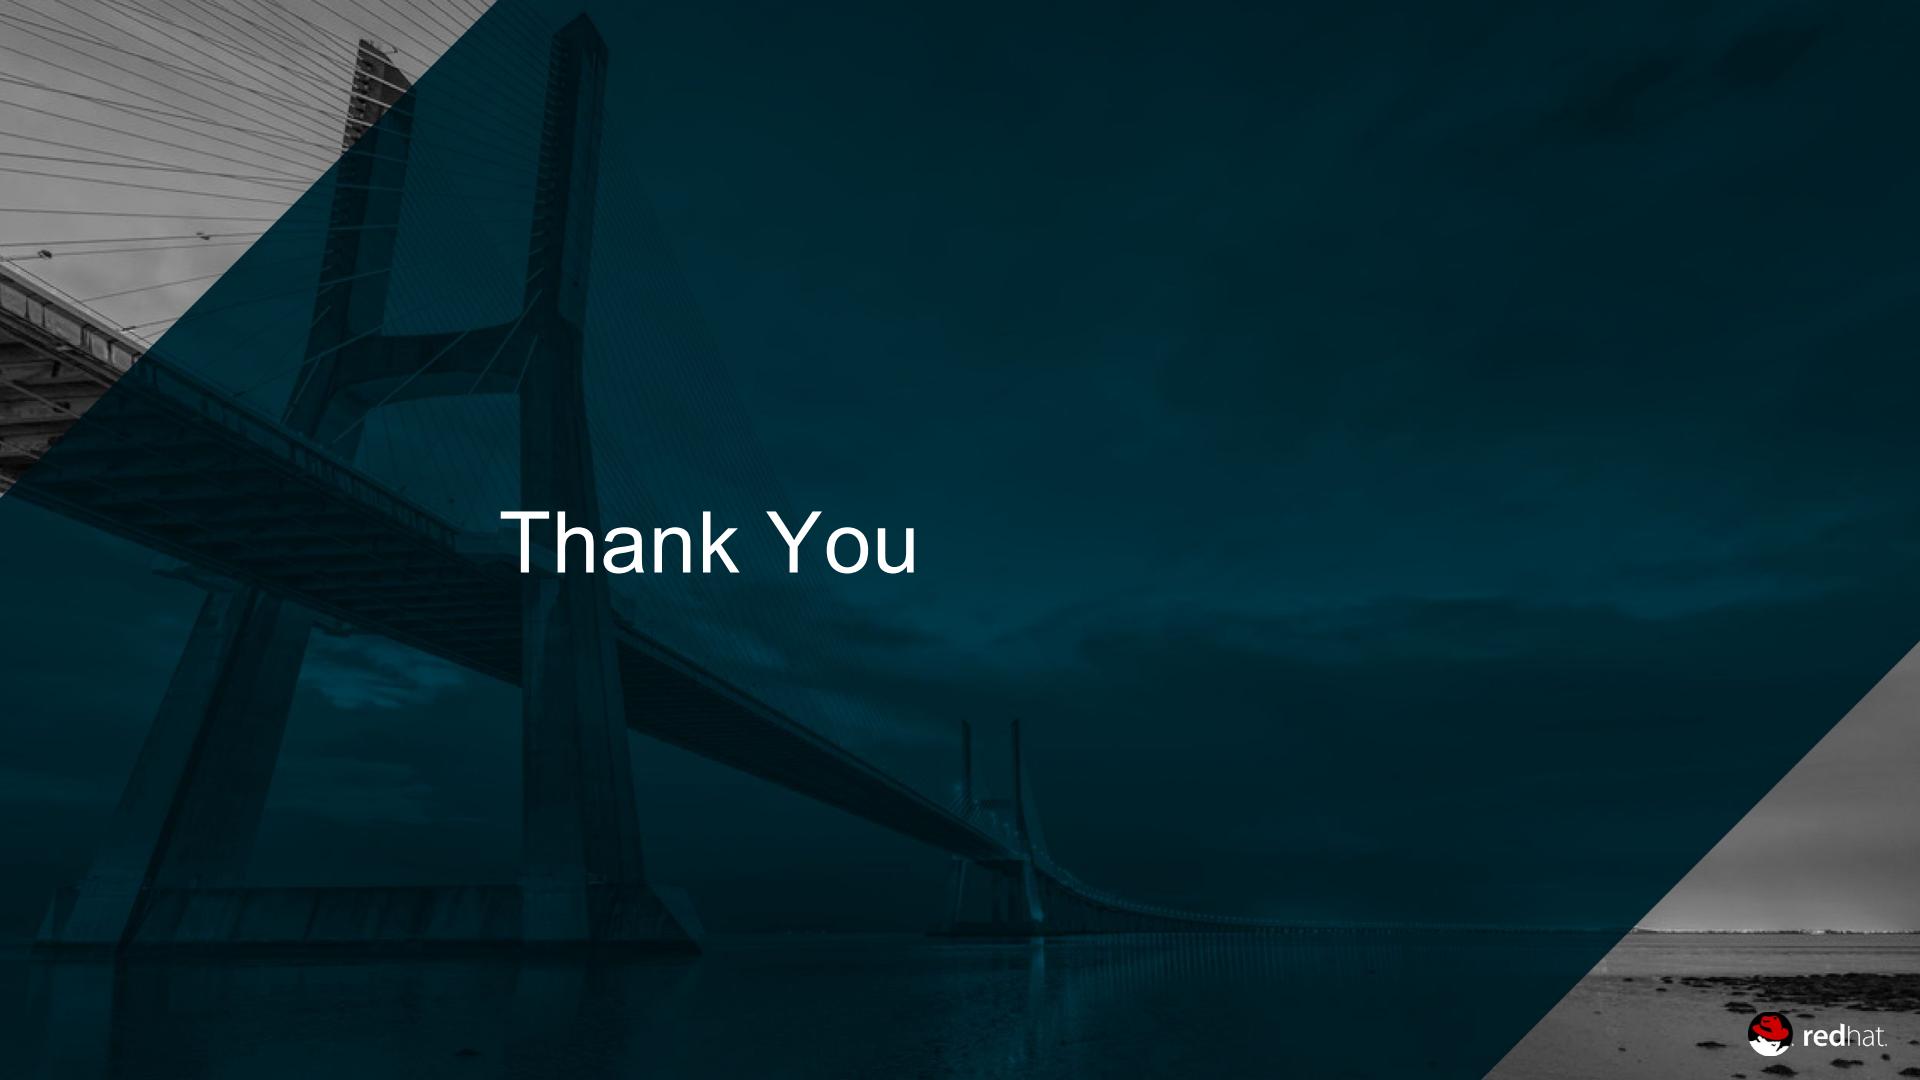
# hello world app - v2

Kubernetes Cluster



# Demo Placeholder

Canary Deployment



# Thank You